

## About:

This document is aimed to provide detail and structured information about NFSv4, ACLs and contains description of the test cases (implemented in Python) and my work about discovery NFSv4 and ACLs.

## Creator:

**A.V.Nesterovich@gmail.com**

## References:

Title	Reference
RFC7530 Network File System (NFS) Version 4 Protocol	<a href="https://www.rfc-editor.org/rfc/rfc7530.txt">https://www.rfc-editor.org/rfc/rfc7530.txt</a>
RFC7531 Network File System (NFS) Version 4 External Data Representation Standard (XDR) Description	<a href="https://www.rfc-editor.org/rfc/rfc7531.txt">https://www.rfc-editor.org/rfc/rfc7531.txt</a>
Mapping Between NFSv4 and Posix Draft ACLs	<a href="http://www.citi.umich.edu/projects/nfsv4/rfc/draft-ietf-nfsv4-acl-mapping-05.txt">http://www.citi.umich.edu/projects/nfsv4/rfc/draft-ietf-nfsv4-acl-mapping-05.txt</a>
acl - Access Control Lists (Linux man page) *POSIX	<a href="http://linux.die.net/man/5/acl">http://linux.die.net/man/5/acl</a>
exports - NFS server export table (Linux man page)	<a href="http://linux.die.net/man/5/exports">http://linux.die.net/man/5/exports</a>
exportfs - maintain table of exported NFS file systems(Linux man page)	<a href="http://linux.die.net/man/8/exportfs">http://linux.die.net/man/8/exportfs</a>
rpc.mountd - NFS mount daemon (Linux man page)	<a href="http://linux.die.net/man/8/mountd">http://linux.die.net/man/8/mountd</a>
getfacl - get file access control lists (Linux man page)	<a href="http://linux.die.net/man/1/getfacl">http://linux.die.net/man/1/getfacl</a>
setfacl - set file access control lists (Linux man page)	<a href="http://linux.die.net/man/1/setfacl">http://linux.die.net/man/1/setfacl</a>


## Info

NFS is a UNIX protocol for large scale client/server file sharing. It is analogous to the server Message Block (SMB) and Common Internet File System (CIFS) protocols on Microsoft Windows. The Network File System Version 4 is a distributed filesystem protocol which owes heritage to NFSv2 and NFSv3. Unlike previous versions of NFS the present version(NFSv4) supports traditional file access while integrating support for file locking and mount protocol.

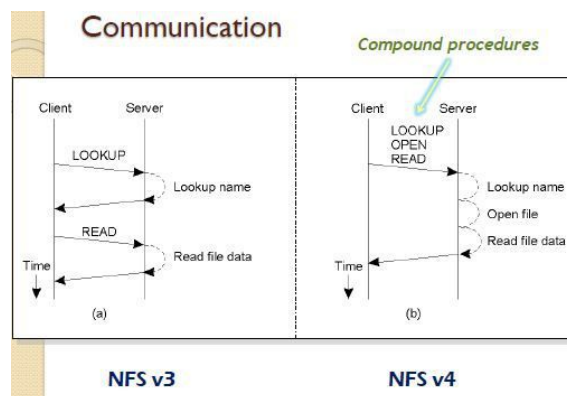
NFSv4 is the successor of NFSv3. It has been designed to work on a LAN or over the Internet.

**NFS v4 provides the following benefits over NFSv3 or earlier NFS versions:**

- **advanced security management** (mandates security and ACLs); Kerberos; SPKM; LIPKEY;
- **firewall friendly** (NFS v4 work by default works over TCP): the use of portmapper dishing out arbitrary ports made it difficult for the firewall. NFSv4 changed that by consolidating most of the TCP/IP services into well-known ports which the security administrator can define in the firewall



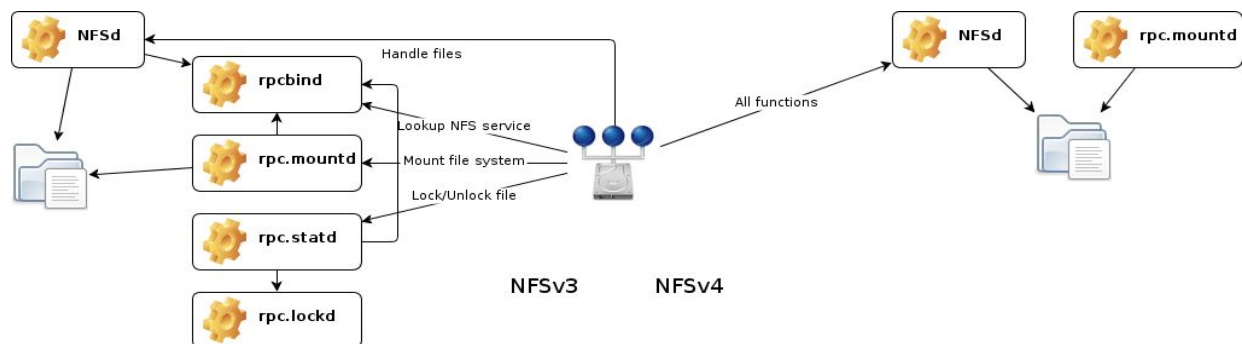
- **advanced and aggressive cache management;**
- **non Unix compatibility (Windows);**
- **easy to administer (Replication, migration);**
- **crash recovery (Client and server sides);**
- **performance improvements:** one key enhancement is the introduction of the COMPOUND RPC procedure which allows the NFS client to group together a bunch of file operations into a single request to the NFS server. This not only reduces the network round-trip latency, but also reduces the small little chatters of the smaller file operations.



- **stateful:** NFSv3 is stateless and it does not maintain the state of the NFS clients. NFSv4 is stateful

and implements a mandatory locking and delegation mechanisms. Leases for locks from the servers to the clients was introduced. A lease is a time-bounded grant for the control of the state of a file and this is implemented through locks.

The NFSv3 and NFSv4 protocols are not compatible. A NFSv4 client cannot access a NFSv3 server, and vice versa. However, in order to simplify migrations from NFSv3 to NFSv4, both NFSv3 and NFSv4 services are launched by the command: **rpc.nfsd**. In the case of NFSv3 and NFSv4 clients simultaneously accessing the same server, one must be aware that two different file systems are used: there is no backward support to NFSv3 by the NFSv4 server. In order to ensure a better reliability over the Internet, NFSv4 only uses TCP. To help NFS setup for internet use, one unique network port is used on NFSv4. This predetermined port is fixed. The default is **port 2049**. With NFSv4, mount and locking services have been integrated in the NFS daemon itself.



Compare NFSv3 and NFSv4 implementations

The NFSv4 server and client work **without the portmap, rpc.lockd, rpc.statd** daemons. The **rpc.mountd daemon is still required** on the server.

Since NFSv4 no longer utilizes the **rpc.mountd** protocol as was used in NFSv2 and NFSv3, the mounting of file systems has changed. An NFSv4 client now has the ability to see all of the exports served by the NFSv4 server as a single file system, called the **NFSv4 pseudo-file system**. On Red Hat Enterprise Linux, the pseudo-file system is identified as a single, real file system, identified at export with the **fsid=0** option.

A NFSv4 client communicates with corresponding NFSv4 Server via Remote Procedure Calls (RPC's). The client sends a request and gets a reply from the server. A NFSv4 server can only provide/export a single, hierarchical file system tree. If a server has to share more than one logical file system tree, the single trees are integrated in a new virtual root directory. This construction, called pseudo file system, is the one which is provided/exported to clients.

# ACL


## General Info

**POSIX** is a family of standards, specified by the IEEE, to clarify and make uniform the application programming interfaces (and ancillary issues, such as commandline shell utilities) provided by Unix-y operating systems.

**ACL** = Access Control List (ACLs allow a sysadmin to express nontrivial rules defining access control on objects (e.g. files or directories))

Option	Access Model In General	Access Models For Filesystems
Subject	entity which performs actions	process
Object	entity on which actions are performed	inode, i.e. file or directory or other object
Algorithm	<b>(subjectcredentials, objectpermissions, actionrequested)</b> 1) definition of format for subject credentials 2) definition of format for object permissions set of actions a subject can perform	<b>might use the permissions of the parent directory</b> 1) subject credentials = an owner, some groups 2) object permissions = an owner, an owning group, a POSIX mode and/or some kind of ACL actions e.g. Read, Write, Execute, Delete, Change Owner

### Access Models For Filesystems:

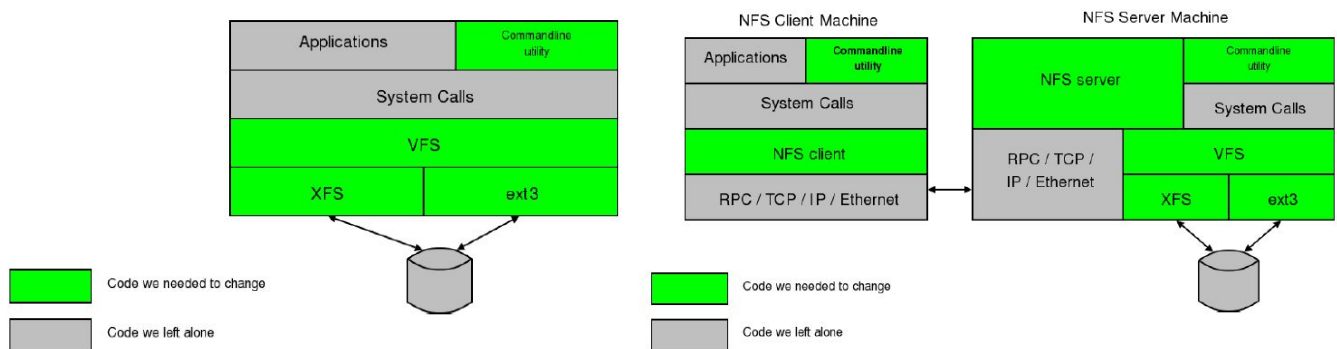


Model	Description
<b>Traditional POSIX</b>	Owners & groups identified by UIDs, GIDs / Locally mapped on each system (e.g. /etc/passwd, NIS (Network Information System, LDAP (Lightweight Directory Access Protocol)) / 3 bit access mask / Read (r), Write (w), Execute (x) / Most actions are mapped to one of these 3 / – exceptions, e.g. Delete => Write on the parent directory / Subjects classified into one of 3 classes / – Owner = the subject's owner matches the object's, or / – Group = one of the subject's groups matches the object's owning group, or / – Other = none of the above / 1 access mask per class = 9 bits / setuid, setgid, sticky bits = 12 bits mode (Total) / Classify subject by matching UIDs and GIDs / - Use class to choose one of the 3 access masks / - If desired bit is set access is allowed, else denied
<b>Draft POSIX ACLs</b>	Draft POSIX extensions 1003.1e and 1003.2c / – never ratified / – but implemented several times

	<ul style="list-style-type: none"> <li>– simply extends POSIX to allow more entries / Users &amp; groups UIDs/GIDs (same) / 3 bit access mask (same) / 3 special classes (same) / – Owner (tag ACL_USER_OBJ), Group (ACL_GROUP_OBJ), Other (ACL_OTHER) / 1 entry per each special class / Any number of additional entries (ACEs)</li> <li>– General user (ACL_USER), general group (ACL_GROUP) / – Total between 3 .. _POSIX_ACL_ENTRIES_MAX (&gt;= 16) entries / – Order not significant / Entries only allow access, never deny access / – Access not explicitly allowed is implicitly denied / If subject UID matches an ACL_USER or ACL_USER_OBJ entry, / use that / Else if subject GIDs match an ACL_GROUP_OBJ or ACL_GROUP entry, use that / Else use the ACL_OTHER entry</li> </ul>
<b>Windows NT ACLs</b>	<p>From Microsoft / – Implemented in Windows NT / – Minor changes in subsequent releases</p> <p>Users &amp; groups identified by SIDs (Security Identifiers) / – Like a variable-length enormous binary UID but with global scope / – e.g. S152110043363481177238915682003330512 / 14 access mask bits / – ReadData/ListFolder, WriteData/CreateFile, AppendData/CreateFolder, /ReadExtendedAttributes, WriteExtendedAttributes, Execute/TraverseFolder, / DeleteChild, ReadAttributes, WriteAttributes, Delete, ReadPermissions, / WritePermissions, TakeOwnership, Synchronize / Variable number of ACEs / – Up to 64K size = ~1800 ACEs / – Entry has a SID, a type, some flags, and an access mask / – Order significant / 3 useful special classes (WellKnown SIDs)</p> <ul style="list-style-type: none"> <li>– Creator Owner, Creator Group, Everyone /Several less helpful WellKnown SIDs / – Interactive, Network, Dialup, Batch, Anonymous, Authenticated, Service etc / Any number of other SIDs</li> <li>– general user, general group / Entry type / – AccessAllowed = subject is allowed the access bits</li> <li>– AccessDenied = subject is denied the access bits / – SystemAudit = wacky audit stuff</li> </ul> <p>Entry flags / – INHERITED, INHERIT_ONLY, CONTAINER_INHERIT, OBJECT_INHERIT, / NO_PROPAGATE_INHERIT – supports inheritance / – SUCCESSFUL_ACCESS, FAILED_ACCESS – wacky audit stuff / ACL flags / – Actually in the containing Security Descriptor / – AUTO_INHERITED, DEFAULTED, PROTECTED / Loop through each entry / – stop if entry's SID matches any of subject's SIDs / If matching entry is Allow, access allowed / If matching entry is Deny, access denied / If no matching entry, access denied / TODO:reexpress</p>
<b>NFSv4 ACLs</b>	<p>Defined in NFSv4 standards / Tries to make the Windows access model usable over NFS / – without admitting it's Windows / – obviously nobody implemented it before writing the RFC / – architecture very similar to Windows, differs only in details / Users &amp; groups identified by strings</p> <ul style="list-style-type: none"> <li>– “user@domain” or “group@domain” / – for transmission; systems expected to map these to some other local form/ like UIDs / 14 access mask bits / – Binary values identical to Windows</li> <li>– Names and semantics...similar...to Windows / – NFSv4.1 adds 2 more which have no equivalent in Windows / 3 useful special classes (whos) / – OWNER@, GROUP@, EVERYONE@</li> <li>– Nearly identical to Windows / Several unhelpful special classes / – Blindly copied from Windows</li> <li>– But much less helpful in an NFS context / Variable number of other entries TODO:rephrases</li> <li>– General user, general group / Any number of entries / – No defined limit / Entry type</li> <li>– Blindly copied from Windows / – Including audit / Per-entry flags / – Blindly copied from Windows / – Added ACE4_IDENTIFIER_GROUP to tell apart user &amp; group whos / Per-ACL flags</li> <li>– Blindly copied from Windows / – Not in NFSv4, added in 4.1</li> </ul>

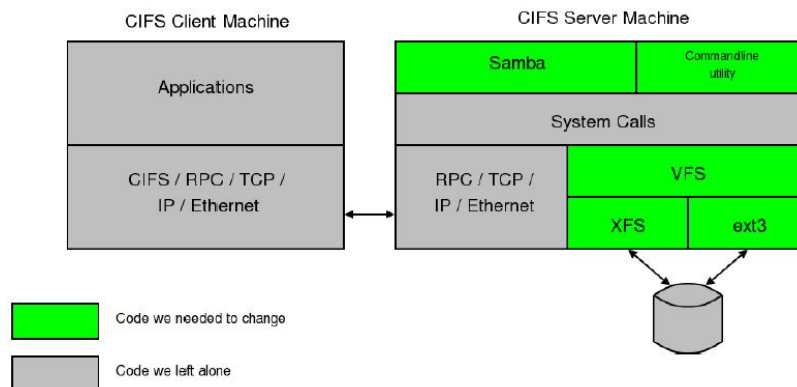
## Implement NFSv4 ACLs in:

NFS4 ACL (Server)	NFS4 ACL (Client)
<p>Server assumes underlying filesystem does POSIX ACLs</p> <ul style="list-style-type: none"> <li>– Converts to POSIX ACL when client sets an ACL</li> <li>– Converts from POSIX ACL when client gets an ACL</li> <li>– In general this conversion is lossy</li> <li>– Samba is doing (hopefully) the same conversion in userspace</li> </ul>	<p>Client presents ACLs in an unexpected manner</p> <ul style="list-style-type: none"> <li>– Nonstandard Extended Attribute, formatted as NFSv4 XDR</li> <li>– Need special utilities to set, print</li> <li>– These utilities are different to what's used on the server</li> <li>– Problems with cp, tar</li> </ul>
<p><b>NFSv4 ACLs are more expressive:</b> – e.g. Deny ACEs, inheritance control / – you might learn to like that extra power /</p> <p><b>Can a mixed Windows/Linux environment:</b> – with global access policies / – and files being shared between both clients /</p> <p>– over both NFS and CIFS protocols</p>	



Architecture: Local Applications

Architecture: Using NFS



Architecture: Using CIFS

## NFSv4 ACLs vs POSIX

### ACE4\_SYNCHRONIZE access bit

- obviously makes no sense on POSIX

### ACE4\_{READ,WRITE}\_NAMED\_ATTR access bits

- make no sense either, but rather less obviously!

### Preserving more obscure corners of POSIX behaviour

– Sticky bit. CAP\_FOWNER, CAP\_CHOWN. Restricted\_chown.

## Doing chmod right: file\_masks and the protocol

### EVERYONE@ != POSIX Other class

– Far too easy to forget

### Other whos (INTERACTIVE@ etc)

– Make no sense in NFS context; preserved but ignored

## ACL Text Representation

List of ACEs, separated by whitespace

Each ACE is four fields separated by colons

who:accessmask:flags:type (e.g. accounts:rwax:g:allow)

#### 1) who

#### 2) access mask – 1char abbrevs, any order

r = read\_data/list\_directory

w = write\_data/add\_file

a = append\_data/add\_subdirectory

x = execute / traverse\_directory,

etc etc

#### 3) Flags – 1char abbrevs, any order

g = who field names a group

#### 4) Type

allow, deny

```
$ ls -l myfile
-rw-rw-r-- 1 me us 0 2008-12-10 10:03 myfile

$ nfs4acl --get myfile
myfile:
owner@:rw::allow
group@:rw::allow
everyone@:r::allow
```

Annotations:

- 'r' = read\_data, 'w' = write\_data
- Annotations for the ACL entries:
  - owner@:rw::allow
    - owner@: who
    - rw: access
    - :: flags
    - allow: type
  - group@:rw::allow
    - group@: who
    - rw: access
    - :: flags
    - allow: type
  - everyone@:r::allow
    - everyone@: who
    - r: access
    - :: flags
    - allow: type

## Test cases (task)

### Intro

Task verifies several candidate abilities including:

- ability to study independently by learning POSIX file systems standard
- perform data analysis by selecting important information
- design test cases and prepare documentation
- implement code based on design

### Test Task

Design a set of test cases for owner / permission / content modification testing of NFSv4 file system.

Implement designed test cases as testing application (test suite). E.g, all tests are stored in “tests” folder and there is “main” file which run all tests and produces an output.

### The results of the task are:

1. Test documentation. Use the following format:

- Name of test case
- Description
- Steps
- Expected result of each step

2. Source code of test suite

3. Logs of the latest successful tests execution

### Test case example:

- Test name: Change file attributes to disable run, enable it, disable again.
- Description: test verifies that after several disable, enable actions permissions set to last value.

### Acceptance criteria:

1. At least 6 test cases have to be created
2. At least 2 test cases for ACL management verification (optional)
3. Test documentation
4. Use one of the following scripting language:
  - a. Python (preferable)
  - b. Ruby
  - c. Perl (in OOP style)
5. Test Suite has to prepare and clean environment
6. Keep logs in log file, Short summary should be printed at the end of testing. E.g.:  
TC001: Passed    TC002: Failed    TC003: Passed
7. Should be executable at any Linux-Like system
8. Please use comments in the code



## Tests cases (Implementation)

### Goal:

*To develop the test suite and create documentation (subject NFSv4 with ACL support test automatization for Linux-like systems [server-client sides])*

### Test environment

Hostname	IP	Software	Description
fedora	192.168.100.182 (LAN)	Fedora 23 Workstation x86-x64 (ext4) <b>rsync</b>	Python 2.7, IDE Pycharm (create tests) Client NFSv4 (run tests)
nestu	Dynamic IP (WAN)	Fedora 23 Workstation x86-x64 (ext4) <b>rsync</b>	Python 2.7, IDE Pycharm (create tests) Client NFSv4 (run tests)
rhel	192.168.100.176 (LAN)	Red Hat Enterprise Linux Server release 7.2 (Maipo) (RHEL) x86-x64 (LVM, XFS)	Server NFSv4 (run tests)

### Daemons NFS

	both sides	server side
<b>user daemons</b>	<b>rpc.idmapd</b> - This process provides NFSv4 client and server upcalls which map between on-the-wire NFSv4 names (which are strings in the form of user@domain) and local UIDs and GIDs. For idmapd to function with NFSv4, the /etc/idmapd.conf must be configured. This service is required for use with NFSv4.	<b>rpc.nfsd</b> - Allows explicit NFS versions and protocols the server advertises to be defined. It works with the Linux kernel to meet the dynamic demands of NFS clients, such as providing server threads each time an NFS client connects. This process corresponds to the nfs service.  <b>rpc.mountd</b> - daemon is still needed to handle the exports, but is not involved with network communication anymore (in other words, the client connects directly with the NFS daemon).
<b>kernel parts</b>	NFSv4, RPC, XDR, TCP, IPv4	

### Packets and tools NFS

Name	Description	Addition
nfs-utils	The nfs-utils package provides a daemon for the kernel NFS server and related tools, which provides a much higher level of performance than the traditional Linux NFS server used by most users.	-
nfs4-acl-tools	The nfs4-acl-tools packages provide utilities for managing NFSv4 Access Control Lists (ACLs) on files and directories mounted on ACL-enabled NFSv4 file systems. These updated packages fix the following bug.	-
libnfsidmap	Is a library holding multiple methods of mapping names to id's and visa versa, mainly for NFSv4.	-
showmount	show mount information for an NFS server	-

## Configs NFS

File	Description	Comment
<b>Server side</b>		
/etc/exports	It is a main configuration file, controls which file systems are exported to remote hosts and specifies options. This file contains a list of entries; each entry indicates a volume that is shared and how it is shared.	There must be at least one entry with fsid=0. (this will be pseudo file system's /). <b>acl</b> option use.
/etc/sysconfig/nfs	This file is used to control which ports the required RPC services run on. Here the number of kernel threads, NFSv4 support and GSS security (kerberos) for NFS can be configured.	Not used on the project
/etc/idmapd.conf	It translates user and group ids into names, and to translate user and group names. Use to modify the default "Domain" to contain DNS domain name.	Not used on the project
<b>Client side</b>		
/etc/fstab	This file is used to control what file systems including NFS directories are mounted when the system boots.	<b>acl</b> option use
/etc/idmapd.conf	It translates user and group ids into names, and to translate user and group names. Use to modify the default "Domain" to contain DNS domain name.	Not used on the project

## Useful commands NFS

Description (Action)	Client side (Comand)	Server side (Comand)
Umount all nfs mounts on client	umount -a -t nfs	-
Check all nfs mounts on client	mount   grep nfs	-
Reexport of shares files	-	exportfs -r
Display of shares files		exportfs -v
Check all registered RPC programs (nfs, portmapper, mountd, ...)	rpcinfo -p	rpcinfo -p
Check mount information on NFS server	showmount -e <server IP>	showmount -e <server IP>
Display statistics kept about NFS client and server activity	nfsstat	nfsstat
Get file access control lists. For each file, getfacl displays the file name, owner, the group, and the Access Control List (ACL). If a directory has a default ACL, getfacl also displays the default ACL. Non-directories cannot have default ACLs.	getfacl -R <dir>	getfacl -R <dir>
Set file access control lists	setfacl <options>	setfacl -R <options>

## Install and settings of NFSv4 (\*\* without Kerberos)

### Server-side (hostname: rhel)

#### 1) Modify hosts file in order to resolve IP from hostname

```
vim /etc/hosts
```

```
192.168.100.176 rhel
```

```
192.168.100.182 fedora
```

```
ping fedora
```

```
PING fedora (192.168.100.182) 56(84) bytes of data.
```

```
64 bytes from fedora (192.168.100.182): icmp_seq=1 ttl=64 time=0.335 ms
```

#### 2) In order to use ACLs enable mount option

```
vim /etc/fstab
```

```
dev/mapper/rhel_nfs-root / xfs defaults,acl 0 0
```

```
mount -a
```

#### 3) Check which loadable kernel modules NFS currently loaded

```
lsmod | grep nfs
```

```
nfs                251815 0
fscache            64987 1 nfs
nfsd               302351 1
auth_rpcgss        59314 1 nfsd
nfs_acl            12837 1 nfsd
lockd              93572 2 nfs,nfsd
grace              13288 2 nfsd,lockd
sunrpc             300421 8 nfs,nfsd,auth_rpcgss,lockd,nfs_acl
```

#### 4) Check and show information about NFS modules:

```
modinfo nfs
```

```
filename: /lib/modules/3.10.0-327.13.1.el7.x86_64/kernel/fs/nfs/nfs.ko
```

```
modinfo nfsv3
```

```
filename: /lib/modules/3.10.0-327.13.1.el7.x86_64/kernel/fs/nfs/nfsv3.ko
```

```
modinfo nfsv4
```

```
filename: /lib/modules/3.10.0-327.13.1.el7.x86_64/kernel/fs/nfs/nfsv4.ko
```

```
modinfo nfsd
```

```
filename: /lib/modules/3.10.0-327.13.1.el7.x86_64/kernel/fs/nfsd/nfsd.ko
```

```
modinfo nfs_acl
```

```
filename: /lib/modules/3.10.0-327.13.1.el7.x86_64/kernel/fs/nfs_common/nfs_acl.ko
```

```
modinfo nfs_layout_flexfiles
```

```
filename: /lib/modules/3.10.0-327.13.1.el7.x86_64/kernel/fs/nfs/flexfilelayout/nfs_layout_flexfiles.ko
```

```
modinfo nfs_layout_nfsv41_files
```

```
filename: /lib/modules/3.10.0-327.13.1.el7.x86_64/kernel/fs/nfs/filelayout/nfs_layout_nfsv41_files.ko
```

#### 5) Check linux kernel for ACL support (find \*=Y option in accordance with the task)

```
uname -a
```

```
Linux rhel 3.10.0-327.13.1.el7.x86_64 #1 SMP Mon Feb 29 13:22:02 EST 2016 x86_64 x86_64 x86_64 GNU/Linux
```

```
grep -i acl /boot/config*
```

```
/boot/config-3.10.0-327.13.1.el7.x86_64:CONFIG_EXT4_FS_POSIX_ACL=y
```

```
/boot/config-3.10.0-327.13.1.el7.x86_64:CONFIG_XFS_POSIX_ACL=y
```

```
/boot/config-3.10.0-327.13.1.el7.x86_64:CONFIG_BTRFS_FS_POSIX_ACL=y
```

```

/boot/config-3.10.0-327.13.1.el7.x86_64:CONFIG_FS_POSIX_ACL=y
/boot/config-3.10.0-327.13.1.el7.x86_64:CONFIG_GENERIC_ACL=y
/boot/config-3.10.0-327.13.1.el7.x86_64:CONFIG_TMPFS_POSIX_ACL=y
/boot/config-3.10.0-327.13.1.el7.x86_64:CONFIG_NFS_V3_ACL=y
/boot/config-3.10.0-327.13.1.el7.x86_64:CONFIG_NFSD_V2_ACL=y
/boot/config-3.10.0-327.13.1.el7.x86_64:CONFIG_NFSD_V3_ACL=y
/boot/config-3.10.0-327.13.1.el7.x86_64:CONFIG_NFS_ACL_SUPPORT=m
/boot/config-3.10.0-327.13.1.el7.x86_64:CONFIG_CIFS_ACL=y

```

If there is N instead of Y, then it means linux kernel doesn't support ACL and need to be recompiled in accordance with the task.

## 6) Install NFS server and tools

```
yum install nfs-utils nfs4-acl-tools libnfsidmap
```

## 7) Check and enable NFS server services

```
systemctl list-unit-files | grep nfs
```

```

proc-fs-nfsd.mount          static
var-lib-nfs-rpc_pipefs.mount static
nfs-blkmap.service          disabled
nfs-config.service          static
nfs-idmap.service           static
nfs-idmapd.service          static
nfs-lock.service            static
nfs-mountd.service          static
nfs-rquotad.service         disabled
nfs-secure-server.service   static
nfs-secure.service          static
nfs-server.service          disabled
nfs-utils.service           static
nfs.service                 disabled
nfslock.service             static
nfs-client.target           enabled

```

```
systemctl enable nfs-server.service
```

```
systemctl enable nfs.service
```

```
systemctl start nfs-server.service
```

```
systemctl start nfs.service
```

## 8) Create NFS share and change permissions (the export filesystem) - a directories to share with client servers

```
mkdir -p /nfs
```

```
chmod a+rwxt /export
```

## 9) Share directories of NFS server for any (LAN, WAN, ...). Exports - NFS server export table.

```
vim /etc/exports
```

```
/nfs *(rw,fsid=0,nohide,no_root_squash,insecure,no_subtree_check,sync)
```

```
/nfs/tests *(rw,nohide,insecure,no_subtree_check,sync)
```

**/nfs** and **/nfs/tests** - shared directories

**\*** - users from any IP address of client machine are allowed to mount directories (means any client)

**rw** - allow both read and write requests on this NFS volume. The default is to disallow any request which changes the filesystem.

**fsid=0** - export a directory over NFSv4. NFSv4 has a concept of a root of the overall exported

filesystem. The export point exported with fsid=0 will be used as this root. The /nfs directory will be root for clients. For example, if you got /nfs/tests subdirectory, then client would see them as /tests directory. NFS needs to be able to identify each filesystem that it exports. For NFSv4, there is a distinguished filesystem which is the root of all exported filesystem. This is specified with fsid=root or fsid=0 both of which mean exactly the same thing. Only for “root” directory.

**nohide** - setting the nohide option on a filesystem causes it not to be hidden, and an appropriately authorised client will be able to move from the parent to that filesystem without noticing the change.

**no\_root\_squash** - turn off root squashing. This option is mainly useful for diskless clients. By default, any file request made by user root on the client machine is treated as by user nobody on the server. (Exactly which UID the request is mapped to depends on the UID of user “nobody” on the server, not the client.) If no\_root\_squash is selected, then root on the client machine will have the same level of access to the files on the system as root on the server.

**insecure** - option in this entry allows clients with NFS implementations that don't use a reserved port for NFS.

**no\_subtree\_check** - this option disables subtree checking, which has mild security implications, but can improve reliability in some circumstances. If a subdirectory of a filesystem is exported, but the whole filesystem isn't then whenever a NFS request arrives, the server must check not only that the accessed file is in the appropriate filesystem (which is easy) but also that it is in the exported tree (which is harder). This check is called the subtree\_check.

**sync** - reply to requests only after the changes have been committed to stable storage (if async - improve performance, but at the cost that an unclean server restart (i.e. a crash) can cause data to be lost or corrupted). All changes to the according filesystem are immediately flushed to disk; the respective write operations are being waited for.

## 10) Restart NFS server services

```
systemctl restart nfs-server.service
```

```
systemctl restart nfs.service
```

## 11) Reexport all directories after modifying /etc/exports and display a list of shares files and export options on a NFS server

```
exportfs -r
```

```
exportfs -v
```

```
/nfs<world>(rw,wdelay,nohide,insecure,no_root_squash,no_subtree_check,fsid=0,sec=sys,rw,insecure,no_root_squash,no_all_squash)
```

```
/nfs/tests<world>(rw,wdelay,nohide,insecure,root_squash,no_subtree_check,fsid=0,sec=sys,rw,insecure,root_squash,no_all_squash)
```

## 12) Config or Disable firewall (firewalld or iptables services) on NFS server to allow client servers to access NFS shares. Open TCP port # 2049 which is used by NFSv4.

### a) Config firewall

```
firewall-cmd --permanent --add-service nfs *** need only one for remote mount via TCP port 2049
```

```
firewall-cmd --permanent --add-service rpc-bind
```

```
firewall-cmd --permanent --add-service mountd
```

```
firewall-cmd --reload
```

```
firewall-cmd --list-all
```

```
public (default, active)
```

```
interfaces: ens192
```

```
sources:
```

```
services: dhcpv6-client mountd nfs rpc-bind ssh
```

```
ports:
```

*masquerade: no*  
*forward-ports:*  
*icmp-blocks:*  
*rich rules:*

*cat /etc/services | grep mountd*

*mountd 20048/tcp # NFS mount protocol*  
*mountd 20048/udp # NFS mount protocol*

*cat /etc/services | grep nfs*

*nfs 2049/tcp nfsd shilp # Network File System*  
*nfs 2049/udp nfsd shilp # Network File System*  
*nfs 2049/sctp nfsd shilp # Network File System*

*cat /etc/services | grep rpcbind*

*sunrpc 111/tcp portmapper rpcbind # RPC 4.0 portmapper TCP*  
*sunrpc 111/udp portmapper rpcbind # RPC 4.0 portmapper UDP*

## **b) Disable firewall**

*systemctl disable firewalld*

*systemctl stop firewalld*

*systemctl status firewalld*

*systemctl status firewalld*

- *firewalld.service - firewalld - dynamic firewall daemon*

*Loaded: loaded (/usr/lib/systemd/system/firewalld.service; disabled; vendor preset: enabled)*

*Active: inactive (dead)*

## **13) Disable and SELinux**

*vim /etc/selinux/config*

*SELINUX=disable*

*\*\*\**

*sestatus*

*SELinux status: disabled*

## **14) Check all registered RPC programs (nfs, portmapper, mountd)**

*rpcinfo -p*

*program vers proto port service*  
*100000 4 tcp 111 portmapper*  
*100000 4 udp 111 portmapper*  
*100024 1 udp 33569 status*  
*100024 1 tcp 41527 status*  
*100005 1 udp 20048 mountd*  
*100005 1 tcp 20048 mountd*  
*100003 4 tcp 2049 nfs*  
*100003 4 udp 2049 nfs*  
*100227 3 tcp 2049 nfs\_acl*  
*100227 3 udp 2049 nfs\_acl*  
*100021 1 udp 39271 nlockmgr*  
*100021 1 tcp 60402 nlockmgr*

15) Set NFSv4 ACLs (add r/w permissions for user: he on directory: /nfs)

```
setfacl -R -m u:he:rwX /nfs
```

16) Check NFSv4 ACLs

```
getfacl -R /nfs
```

*getfacl: Removing leading '/' from absolute path names*

**# file: nfs**

*# owner: root*

*# group: root*

*user::rwX*

**user:he:rwX**

*group::r-x*

*mask::rwX*

*other::r-x*

**# file: nfs/tests**

*# owner: root*

*# group: root*

*user::rwX*

**user:he:rwX**

*group::r-x*

*mask::rwX*

*other::r-x*

**# file: nfs/tests/music\_for\_programming\_00-manifesto.mp3**

*# owner: he*

*# group: he*

*user::rwX*

**user:he:rwX**

*group::rwX*

*mask::rwX*

*other::rwX*

**# file: nfs/DDNTestTaskE1-E2-true.pdf**

*# owner: he*

*# group: he*

*user::rwX*

**user:he:rwX**

*group::rwX*

*mask::rwX*

*other::r-x*

## 17) Install and settings addition software need for run testcases

**a) rsh - remote shell access (need in order to receive commands from remote client)**

**Install rsh and rshd:**

```
yum install rsh rsh-server
rpm -qa | grep rsh
rsh-server-0.17-76.el7_1.1.x86_64
rsh-0.17-76.el7_1.1.x86_64
```

**Start rsh-server daemons:**

```
systemctl enable rsh.socket
systemctl enable rlogin.socket
systemctl enable rexec.socket
```

```
systemctl start rsh.socket
systemctl start rlogin.socket
systemctl start rexec.socket
```

*systemctl status rsh.socket*

- *rsh.socket - Remote Shell Facilities Activation Socket*  
*Loaded: loaded (/usr/lib/systemd/system/rsh.socket; enabled; vendor preset: disabled)*  
*Active: active (listening)*

*systemctl status rlogin.socket*

- *rlogin.socket - Remote Login Facilities Activation Socket*  
*Loaded: loaded (/usr/lib/systemd/system/rlogin.socket; enabled; vendor preset: disabled)*  
*Active: active (listening)*

*systemctl status rexec.socket*

- *rexec.socket - Remote Execution Facilities Activation Socket*  
*Loaded: loaded (/usr/lib/systemd/system/rexec.socket; enabled; vendor preset: disabled)*  
*Active: active (listening)*

**Configure rsh-server:**

```
vim /root/.rhosts - allow the user root on the client fedora to log in as root on the target (server)
fedora root
```

*vim /etc/securetty* - enable external root user to execute the command (lists terminals from which root can log in)

```
rsh
rexec
rlogin
```



## Client-side (hostname: fedora)

### 1) Modify hosts file in order to resolve IP from hostname

*vim /etc/hosts*

*192.168.100.176 rhel*

*192.168.100.182 fedora*

*ping rhel*

*PING rhel (192.168.100.176) 56(84) bytes of data.*

*64 bytes from rhel (192.168.100.176): icmp\_seq=1 ttl=64 time=0.594 ms*

*\*\*\**

### 2) Disable firewall

*systemctl disable firewalld*

*systemctl stop firewalld*

*systemctl status firewalld*

*systemctl status firewalld*

• *firewalld.service - firewalld - dynamic firewall daemon*

*Loaded: loaded (/usr/lib/systemd/system/firewalld.service; disabled; vendor preset: enabled)*

*Active: inactive (dead)*

### 3) Disable and SELinux

*vim /etc/selinux/config*

*SELINUX=disable*

*\*\*\**

*sestatus*

*SELinux status: disabled*

### 4) Install nfs utils for client NFS service

*yum install nfs-utils libnfsidmap*

### 5) Enable and start NFS service

*systemctl enable nfs-client.target*

*systemctl start nfs-client.target*

### 6) Check nfs shares to the clients on NFS server

*showmount -e 192.168.100.176*

*Export list for 192.168.100.176:*

*/nfs/tests \**

*/nfs \**

### 7) Mount the exported file system

*mount -t nfs4 192.168.100.176:/ /nfs*

*Observe that only "/" is given instead of the actual exported path name*

## 8) Check mounted NFS file system

*mount | grep nfs*

*nfsd on /proc/fs/nfsd type nfsd (rw,relatime)*

*sunrpc on /var/lib/nfs/rpc\_pipefs type rpc\_pipefs (rw,relatime)*

*192.168.100.176:/ on /nfs type nfs4*

*(rw,relatime,vers=4.1,rsize=524288,wsiz=524288,namlen=255,hard,proto=tcp,timeo=600,retrans=2,sec=sys,clientaddr=192.168.100.182,local\_lock=none,addr=192.168.100.176) \*\*\* ver 4.1*

*df -hT | grep nfs*

*192.168.100.176:/ nfs4 37G 5.5G 32G 15% /nfs*

## 9) Mount NFS file system permanently in order to mount after reboot system

*vim /etc/fstab*

*192.168.100.176:/ /nfs nfs4\_netdev,auto 0 0*

*mount -a*

## 10) Check all registered RPC programs

*sdrpcinfo -p*

*program vers proto port service*

*100000 4 tcp 111 portmapper*

*100000 4 udp 111 portmapper*

## 11) Install and settings additional software need for run testcases

### a) rsh - remote shell access (need in order to execute commands on remote server)

**Install rsh and rshd:**

*yum install rsh rsh-server*

*rpm -qa | grep rsh*

*rsh-server-0.17-76.el7\_1.1.x86\_64*

*rsh-0.17-76.el7\_1.1.x86\_64*





Test #1 <Test of ...>

Test #2 <Test of ...>

Test #3 <Test of ...>

Test #4 <Test of ...>

Test #5 <Test of ...>

## Add

1)

ltp/include/mk/env\_pre.mk

ltp/include/mk/env\_pre.mk

To >>>>>

/cloud/Dropbox/sync/git/python/ltp/testcases/network/nfsv4/acl

2) file:///python/ltp-master/include/mk/env\_pre.mk

find . -type d -exec chmod 755 {} \;

find . -type f -exec chmod 644 {} \;