

## About:

This document is aimed to provide detail and structured information about NFSv4, ACLs and contains description of the test cases (implemented in Python) and my work about discovery NFSv4 and ACLs.

## Creator:

**A.V.Nesterovich@gmail.com**

## References:

Title	Reference
RFC7530 Network File System (NFS) Version 4 Protocol	<a href="https://www.rfc-editor.org/rfc/rfc7530.txt">https://www.rfc-editor.org/rfc/rfc7530.txt</a>
RFC7531 Network File System (NFS) Version 4 External Data Representation Standard (XDR) Description	<a href="https://www.rfc-editor.org/rfc/rfc7531.txt">https://www.rfc-editor.org/rfc/rfc7531.txt</a>
File System Extended Attributes in NFSv4 draft-ietf-nfsv4-xattrs-02	<a href="https://tools.ietf.org/id/draft-ietf-nfsv4-xattrs-02.txt">https://tools.ietf.org/id/draft-ietf-nfsv4-xattrs-02.txt</a>
Mapping Between NFSv4 and Posix Draft ACLs	<a href="http://www.citi.umich.edu/projects/nfsv4/rfc/draft-ietf-nfsv4-acl-mapping-05.txt">http://www.citi.umich.edu/projects/nfsv4/rfc/draft-ietf-nfsv4-acl-mapping-05.txt</a>
acl - Access Control Lists (Linux man page) *POSIX	<a href="http://linux.die.net/man/5/acl">http://linux.die.net/man/5/acl</a>
exports - NFS server export table (Linux man page)	<a href="http://linux.die.net/man/5/exports">http://linux.die.net/man/5/exports</a>
exportfs - maintain table of exported NFS file systems(Linux man page)	<a href="http://linux.die.net/man/8/exportfs">http://linux.die.net/man/8/exportfs</a>
rpc.mountd - NFS mount daemon (Linux man page)	<a href="http://linux.die.net/man/8/mountd">http://linux.die.net/man/8/mountd</a>
getfacl - get file access control lists (Linux man page)	<a href="http://linux.die.net/man/1/getfacl">http://linux.die.net/man/1/getfacl</a>
setfacl - set file access control lists (Linux man page)	<a href="http://linux.die.net/man/1/setfacl">http://linux.die.net/man/1/setfacl</a>
get NFSv4 file/directory access control lists (Linux man page)	<a href="http://linux.die.net/man/1/nfs4_getfacl">http://linux.die.net/man/1/nfs4_getfacl</a>
set NFSv4 file/directory access control lists (Linux man page)	<a href="http://linux.die.net/man/1/nfs4_setfacl">http://linux.die.net/man/1/nfs4_setfacl</a>
maximum number of ACL's available on a directory	<a href="https://access.redhat.com/solutions/68429">https://access.redhat.com/solutions/68429</a>

## Chapter #1 - NFS4


### Info

NFS is a UNIX protocol for large scale client/server file sharing. It is analogous to the server Message Block (SMB) and Common Internet File System (CIFS) protocols on Microsoft Windows. The Network File System Version 4 is a distributed filesystem protocol which owes heritage to NFSv2 and NFSv3. Unlike previous versions of NFS the present version(NFSv4) supports traditional file access while integrating support for file locking and mount protocol.

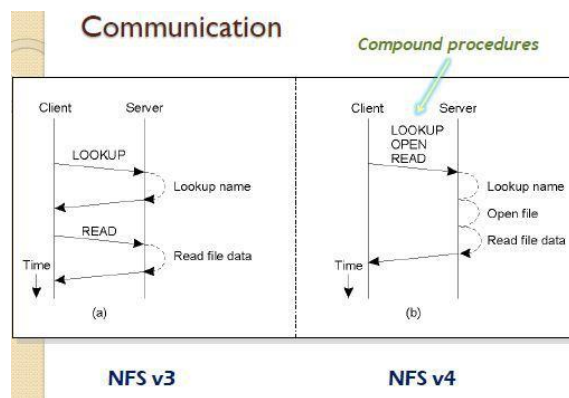
NFSv4 is the successor of NFSv3. It has been designed to work on a LAN or over the Internet.

**NFS v4 provides the following benefits over NFSv3 or earlier NFS versions:**

- **advanced security management** (mandates security and ACLs); Kerberos; SPKM; LIPKEY;
- **firewall friendly** (NFS v4 work by default works over TCP): the use of portmapper dishing out arbitrary ports made it difficult for the firewall. NFSv4 changed that by consolidating most of the TCP/IP services into well-known ports which the security administrator can define in the firewall

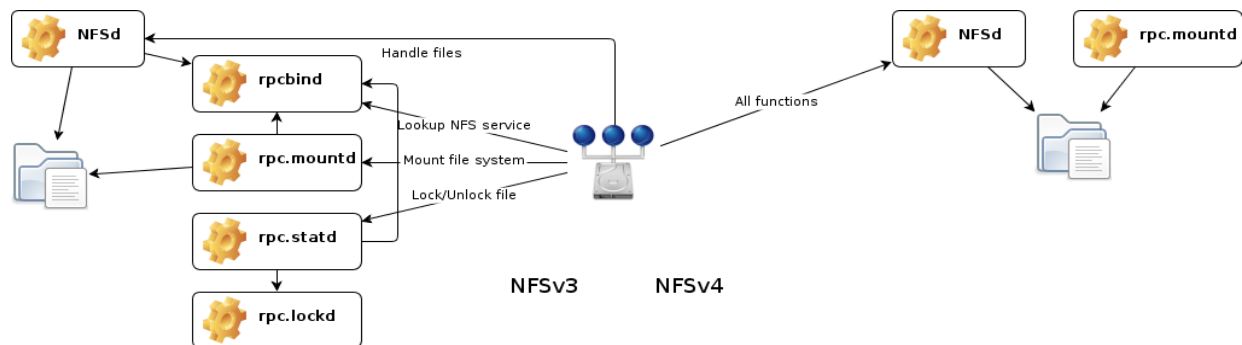


- **advanced and aggressive cache management;**
- **non Unix compatibility (Windows);**
- **easy to administer (Replication, migration);**
- **crash recovery (Client and server sides);**
- **performance improvements:** one key enhancement is the introduction of the COMPOUND RPC procedure which allows the NFS client to group together a bunch of file operations into a single request to the NFS server. This not only reduces the network round-trip latency, but also reduces the small little chatters of the smaller file operations.



- **stateful:** NFSv3 is stateless and it does not maintain the state of the NFS clients. NFSv4 is stateful and implements a mandatory locking and delegation mechanisms. Leases for locks from the servers to the clients was introduced. A lease is a time-bounded grant for the control of the state of a file and this is implemented through locks.

The NFSv3 and NFSv4 protocols are not compatible. A NFSv4 client cannot access a NFSv3 server, and vice versa. However, in order to simplify migrations from NFSv3 to NFSv4, both NFSv3 and NFSv4 services are launched by the command: **rpc.nfsd**. In the case of NFSv3 and NFSv4 clients simultaneously accessing the same server, one must be aware that two different file systems are used: there is no backward support to NFSv3 by the NFSv4 server. In order to ensure a better reliability over the Internet, NFSv4 only uses TCP. To help NFS setup for internet use, one unique network port is used on NFSv4. This predetermined port is fixed. The default is **port 2049**. With NFSv4, mount and locking services have been integrated in the NFS daemon itself.



### Compare NFSv3 and NFSv4 implementations

The NFSv4 server and client work **without the portmap, rpc.lockd, rpc.statd** daemons. The **rpc.mountd daemon is still required** on the server.

Since NFSv4 no longer utilizes the **rpc.mountd** protocol as was used in NFSv2 and NFSv3, the mounting of file systems has changed. An NFSv4 client now has the ability to see all of the exports served by the NFSv4 server as a single file system, called the **NFSv4 pseudo-file system**. On Red Hat Enterprise Linux, the pseudo-file system is identified as a single, real file system, identified at export with the **fsid=0** option.

A NFSv4 client communicates with corresponding NFSv4 Server via Remote Procedure Calls (RPC's). The client sends a request and gets a reply from the server. A NFSv4 server can only provide/export a single, hierarchical file system tree. If a server has to share more than one logical file system tree, the single trees are integrated in a new virtual root directory. This construction, called pseudo file system, is the one which is provided/exported to clients.

## Chapter #2 - ACLs


### General Info

**POSIX** is a family of standards, specified by the IEEE, to clarify and make uniform the application programming interfaces (and ancillary issues, such as commandline shell utilities) provided by Unix-y operating systems.

**ACL** = Access Control List (ACLs allow a sysadmin to express nontrivial rules defining access control on objects (e.g. files or directories))

Option	Access Model In General	Access Models For Filesystems
Subject	entity which performs actions	process
Object	entity on which actions are performed	inode, i.e. file or directory or other object
Algorithm	<b>(subjectcredentials, objectpermissions, actionrequested)</b> 1) definition of format for subject credentials 2) definition of format for object permissions set of actions a subject can perform	<b>might use the permissions of the parent directory</b> 1) subject credentials = an owner, some groups 2) object permissions = an owner, an owning group, a POSIX mode and/or some kind of ACL actions e.g. Read, Write, Execute, Delete, Change Owner

### Access Models For Filesystems:



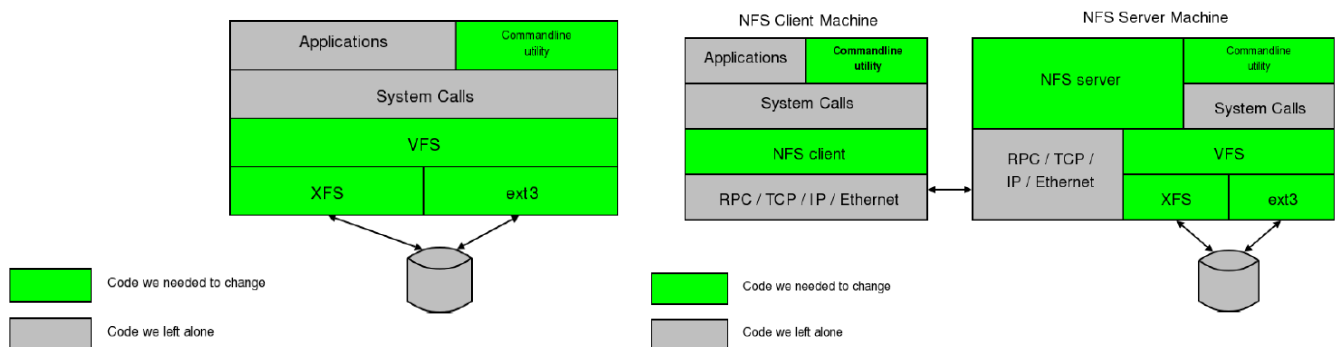
Model	Description
<b>Traditional POSIX</b> <b>chown</b> <b>chmod</b>	Owners & groups identified by UIDs, GIDs / Locally mapped on each system (e.g. /etc/passwd, NIS (Network Information System, LDAP (Lightweight Directory Access Protocol)) / 3 bit access mask / Read (r), Write (w), Execute (x) / Most actions are mapped to one of these 3 / – exceptions, e.g. Delete => Write on the parent directory / Subjects classified into one of 3 classes / – Owner = the subject's owner matches the object's, or / – Group = one of the subject's groups matches the object's owning group, or / – Other = none of the above / 1 access mask per class = 9 bits / setuid, setgid,

	<p>sticky bits = 12 bits mode (Total) / Classify subject by matching UIDs and GIDs / - Use class to choose one of the 3 access masks / - If desired bit is set access is allowed, else denied</p>
<p><b>Draft POSIX ACLs</b></p>	<p>Draft POSIX extensions 1003.1e and 1003.2c / - never ratified / - but implemented several times</p> <ul style="list-style-type: none"> <li>- simply extends POSIX to allow more entries / Users &amp; groups UIDs/GIDs (same) / 3 bit access mask (same) / 3 special classes (same) / - Owner (tag ACL_USER_OBJ), Group (ACL_GROUP_OBJ), Other (ACL_OTHER) / 1 entry per each special class / Any number of additional entries (ACEs) <b>ACE (Access Control Entries)</b> - General user (ACL_USER), general group (ACL_GROUP) / - Total between 3 .. _POSIX_ACL_ENTRIES_MAX (&gt;= 16) entries / - Order not significant / Entries only allow access, never deny access / - Access not explicitly allowed is implicitly denied / If subject UID matches an ACL_USER or ACL_USER_OBJ entry, / use that / Else if subject GIDs match an ACL_GROUP_OBJ or ACL_GROUP entry, use that / Else use the ACL_OTHER entry</li> </ul>
<p><b>Windows NT ACLs</b></p> <p>setfacl</p> <p>getfacl</p>	<p>From Microsoft / - Implemented in Windows NT / - Minor changes in subsequent releases</p> <p>Users &amp; groups identified by SIDs (Security Identifiers) / - Like a variable-length enormous binary UID but with global scope / - e.g. S152110043363481177238915682003330512 / 14 access mask bits</p> <ul style="list-style-type: none"> <li>/ - ReadData/ListFolder, WriteData/CreateFile, AppendData/CreateFolder, /ReadExtendedAttributes, WriteExtendedAttributes, Execute/TraverseFolder, / DeleteChild, ReadAttributes, WriteAttributes, Delete, ReadPermissions, / WritePermissions, TakeOwnership, Synchronize / Variable number of <b>ACEs</b> / - Up to 64K size = ~1800 ACEs / - Entry has a SID, a type, some flags, and an access mask / - Order significant / 3 useful special classes (WellKnown SIDs) <ul style="list-style-type: none"> <li>- Creator Owner, Creator Group, Everyone /Several less helpful WellKnown SIDs / - Interactive, Network, Dialup, Batch, Anonymous, Authenticated, Service etc / Any number of other SIDs</li> <li>- general user, general group / Entry type / - AccessAllowed = subject is allowed the access bits</li> <li>- AccessDenied = subject is denied the access bits / - SystemAudit = wacky audit stuff</li> </ul> </li> <li>Entry flags / - INHERITED, INHERIT_ONLY, CONTAINER_INHERIT, OBJECT_INHERIT, / NO_PROPAGATE_INHERIT - supports inheritance / - SUCCESSFUL_ACCESS, FAILED_ACCESS - wacky audit stuff / ACL flags / - Actually in the containing Security Descriptor / - AUTO_INHERITED, DEFAULTED, PROTECTED / Loop through each entry / - stop if entry's SID matches any of subject's SIDs / If matching entry is Allow, access allowed / If matching entry is Deny, access denied / If no matching entry, access denied / TODO:reexpress</li> </ul>
<p><b>NFSv4 ACLs</b></p> <p>nfs4_getfacl</p> <p>nfs4_setfacl</p>	<p>Defined in NFSv4 standards / Tries to make the Windows access model usable over NFS / - without admitting it's Windows / - obviously nobody implemented it before writing the RFC / - architecture very similar to Windows, differs only in details / Users &amp; groups identified by strings</p> <ul style="list-style-type: none"> <li>- “user@domain” or “group@domain” / - for transmission; systems expected to map these to some other local form/ like UIDs / 14 access mask bits / - Binary values identical to Windows</li> <li>- Names and semantics...similar...to Windows / - NFSv4.1 adds 2 more which have no equivalent in Windows / 3 useful special classes (whos) / - OWNER@, GROUP@, EVERYONE@</li> <li>- Nearly identical to Windows / Several unhelpful special classes / - Blindly copied from Windows</li> <li>- But much less helpful in an NFS context / Variable number of other entries TODO:rephrases</li> <li>- General user, general group / Any number of entries / - No defined limit / Entry type</li> </ul>

- Blindly copied from Windows / – Including audit / Per-entry flags / – Blindly copied from Windows / – Added ACE4\_IDENTIFIER\_GROUP to tell apart user & group whos / Per-ACL flags
- Blindly copied from Windows / – Not in NFSv4, added in 4.1

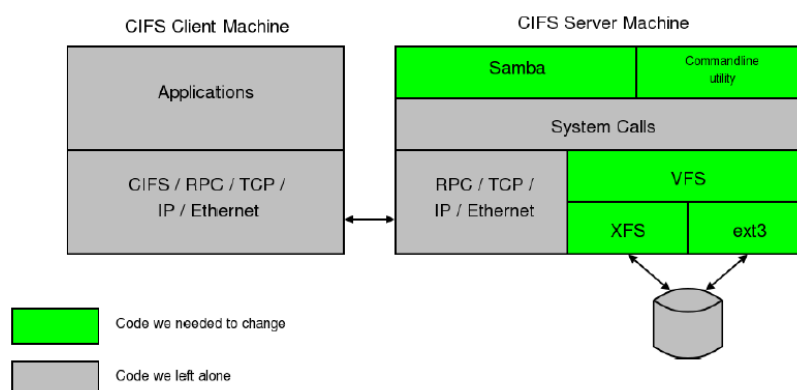
### Implement NFSv4 ACLs in:

NFS4 ACL (Server)	NFS4 ACL (Client)
Server assumes underlying filesystem does POSIX ACLs <ul style="list-style-type: none"> <li>– Converts to POSIX ACL when client sets an ACL</li> <li>– Converts from POSIX ACL when client gets an ACL</li> <li>– In general this conversion is lossy</li> <li>– Samba is doing (hopefully) the same conversion in userspace</li> </ul>	Client presents ACLs in an unexpected manner <ul style="list-style-type: none"> <li>– Nonstandard Extended Attribute, formatted as NFSv4 XDR</li> <li>– Need special utilities to set, print</li> <li>– These utilities are different to what's used on the server</li> <li>– Problems with cp, tar</li> </ul>
<b>NFSv4 ACLs are more expressive:</b> – e.g. Deny ACEs, inheritance control / – you might learn to like that extra power / <b>Can a mixed Windows/Linux environment:</b> – with global access policies / – and files being shared between both clients / – over both NFS and CIFS protocols	



Architecture: Local Applications

Architecture: Using NFS



Architecture: Using CIFS

### NFSv4 ACLs vs POSIX

### **ACE4\_SYNCHRONIZE access bit**

- obviously makes no sense on POSIX

### **ACE4\_{READ,WRITE}\_NAMED\_ATTR access bits**

- make no sense either, but rather less obviously!

### **Preserving more obscure corners of POSIX behaviour**

- Sticky bit. CAP\_FOWNER, CAP\_CHOWN. Restricted\_chown.

### **Doing chmod right: file\_masks and the protocol**

### **EVERYONE@ != POSIX Other class**

- Far too easy to forget

### **Other whos (INTERACTIVE@ etc)**

- Make no sense in NFS context; preserved but ignored

## **ACL Text Representation**

### **List of ACEs, separated by whitespace**

### **Each ACE is four fields separated by colons**

who:accessmask:flags:type (e.g. accounts:rwax:g:allow)

#### **1) who**

#### **2) access mask – 1char abbrevs, any order**

r = read\_data/list\_directory

w = write\_data/add\_file

a = append\_data/add\_subdirectory

x = execute / traverse\_directory,

etc etc

#### **3) Flags – 1char abbrevs, any order**

g = who field names a group

#### **4) Type**



allow, deny

```
$ ls -l myfile
-rw-rw-r-- 1 me us 0 2008-12-10 10:03 myfile

$ nfs4acl --get myfile
myfile:
owner@:rw::allow
group@:rw::allow
everyone@:r::allow
```

who access flags type

## POSIX and ACLs (Windows)

**chown** is for POSIX ACLs permissions (Unix)

**setfacl** is for ACLs (Windows) **ACEs** - NFS share permissions are governed from NFS server side

On NFS server, if file system which is exported by NSF server supports ACL and ACLs **can be read by NFS Clients**, then ACLs are utilized by client System.

For disabling ACLs on NFS share, you have to add option “no\_acl” in ‘/etc/exportfs’ file on NFS Server. To disable it on NSF client side again use “no\_acl” option during mount time.

**Access ACLs:** Access ACLs are used for granting permissions on any file or directory.

**Default ACLs:** Default ACLs are used for granting/setting access control list on a specific directory only.

\*\*\*By default, if the file system being exported by an NFSv4 server supports ACLs and the NFS client can read ACLs, ACLs are utilized by the client system.

## Maximum number of ACL's available on a directory on various filesystems (Supported ACL Entries) (ext2/ext3/ext4; xfs; gfs2; nfs)

The maximum number of ACL's supported on GFS2 and XFS filesystems is hard-coded at 25. This value cannot be altered without changing source code and recompiling the corresponding kernel modules.

The maximum number of ACL's on EXT2/EXT3 on RHEL3 was 32 per file/directory, but this limit was raised in RHEL4 to the maximum number that will fit in a filesystem block.

File system / Info	Restrictions
<b>GFS2</b> The maximum number of ACL's on an individual file/directory on a <b>GFS2 filesystem</b>	<b>25</b> fs/gfs2/acl.h: #define GFS2_ACL_MAX_ENTRIES 25
<b>XFS</b> The maximum number of ACL's on an file/individual directory on an XFS filesystem	<b>25</b> fs/xfs/xfs_acl.h: #define XFS_ACL_MAX_ENTRIES 25
<b>EXT2, EXT3, EXT4</b> The maximum number of ACL's on an individual file/directory on an EXT2, EXT3 or EXT4 filesystem varies with block size	<b>32</b> Based on a blocksize of 4096 (default), approximately 500 ACL's can be stored on each file/directory.  *** Once the limit is reached, attempting to add additional ACL's will result in the error: setfacl: /acltest/directory: No space left on device
<b>NFS</b> The maximum number of ACL's supported on a file/directory exported by NFS	<b>1024</b> include/linux/nfsacl.h: #define NFS_ACL_MAX_ENTRIES 1024

The access ACL of a file system object is accessed for every access decision that involves that object. Access checking is performed on the whole path from the namespace root to the file in question. It is important that ACL access checks are efficient. To avoid frequently looking up ACL attributes and converting them from the machine-independent attribute representation to a machine-specific representation, the Ext2, Ext3, JFS, and ReiserFS implementations cache the machine-specific ACL representations. This is done in addition to the normal file system caching mechanisms, which use either the page cache, the buffer cache, or both. XFS does not use this additional layer of caching.

Most UNIX-like systems that support ACLs limit the number of ACL entries allowed to some reasonable number. ACLs with a high number of ACL entries tend to become more difficult to manage. More than a handful of ACL entries are usually an indication of bad application design. In most such cases, it makes more sense to make better use of groups instead of bloating ACLs.

The ReiserFS and JFS implementations define no limit on the number of ACL entries, so a limit is only imposed by the maximum size of EA values. The current EA size limit is 64 KiB, or 8191 ACL entries, which is too high for ACLs in practice: besides being impractical to work with, the time it would take to check access in such huge ACLs may be prohibitive.

## POSIX Permissions (chown, chgrp, chmod, umask)

### 1) chown : change file owner and group

To check the ownership of a file or directory use `ls -l`

Usage: `chown [-Rcfv] newowner filenames/directory`. Take note only root can change the ownership.

Example:

```
chown linda file.txt
```

This will cause file.txt to now be owned by linda.

```
chown -R abu:sales /home/account/
```

This is going to make all files inside `/home/account/` and its subdirectories to belong to abu and to be associated with the group sales. `-R` means include all subdirectories

### 2) chgrp : change group ownership

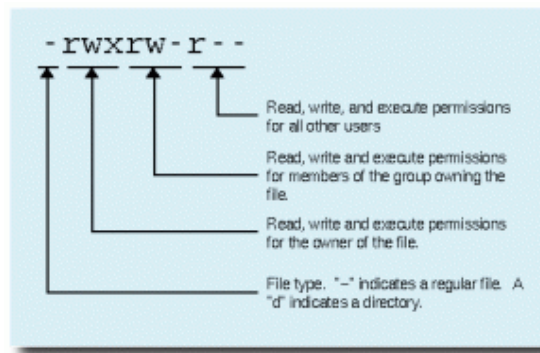
Usage : `chgrp [-Rcfv] groupname foo.txt`

`chgrp marketing file.txt` – to change the group specified to a certain document

`chgrp oracle /usr/database` – to change the group specified to a certain directory

`chgrp -R marketing /sales/2008` – to change the group specified to a certain directory recursively

### 3) chmod : to change the permissions of a file or directory. Use `ls -l` to see the permission settings.



`rwX rwX rwX = 111 111 111`

`rw- rw- rw- = 110 110 110`

`rwX --- --- = 111 000 000`

and so on...

`rwX = 111 in binary = 7`

`rw- = 110 in binary = 6`

`r-x = 101 in binary = 5`

`r-- = 100 in binary = 4`

For example, if we wanted to set `some_file` to have read and write permission for the owner, but wanted to keep the file private from others, we would:

`chmod 600 some_file`

Here is a table of numbers that covers all the common settings. The ones beginning with “7” are used with programs (since they enable execution) and the rest are for other kinds of files.

Value	Meaning
777 (rwxrwxrwx) No restrictions on permissions. Anybody may do anything. Generally not a desirable setting.	777 (rwxrwxrwx) No restrictions on permissions. Anybody may do anything. Generally not a desirable setting.
755 (rwxr-xr-x) The file’s owner may read, write, and execute the file. All others may read and execute the file. This setting is common for programs that are used by all users.	755 (rwxr-xr-x) The file’s owner may read, write, and execute the file. All others may read and execute the file. This setting is common for programs that are used by all users.
700 (rwx—) The file’s owner may read, write, and execute the file. Nobody else has any rights. This setting is useful for programs that only the owner may use and must be kept private from others.	700 (rwx—) The file’s owner may read, write, and execute the file. Nobody else has any rights. This setting is useful for programs that only the owner may use and must be kept private from others.
666 (rw-rw-rw-) All users may read and write the file.	666 (rw-rw-rw-) All users may read and write the file.
644 (rw-r--r-) The owner may read and write a file, while all others may only read the file. A common setting for data files that everybody may read, but only the owner may change.	644 (rw-r--r-) The owner may read and write a file, while all others may only read the file. A common setting for data files that everybody may read, but only the owner may change.
600 (rw—) The owner may read and write a file. All others have no rights. A common setting for data files that the owner wants to keep private.	600 (rw—) The owner may read and write a file. All others have no rights. A common setting for data files that the owner wants to keep private.

Here are some useful settings for directories:

Value	Meaning
777 (rwxrwxrwx) No restrictions on permissions. Anybody may list files, create new files in the directory and delete files in the directory. Generally not a good setting.	777 (rwxrwxrwx) No restrictions on permissions. Anybody may list files, create new files in the directory and delete files in the directory. Generally not a good setting.
755 (rwxr-xr-x) The directory owner has full access. All others may list the directory, but cannot create files nor delete them. This setting is common for directories that you wish to share with other users.	755 (rwxr-xr-x) The directory owner has full access. All others may list the directory, but cannot create files nor delete them. This setting is common for directories that you wish to share with other users.
700 (rwx——) The directory owner has full access. Nobody else has any rights. This setting is useful for directories that only the owner may use and must be kept private from others.	700 (rwx——) The directory owner has full access. Nobody else has any rights. This setting is useful for directories that only the owner may use and must be kept private from others.

## Who

The “who” is a list of letters that specifies whom you’re going to be giving permissions to. These may be specified in any order.

u	The user who owns the file (this means “you.”)
g	The group the file belongs to.
o	The other users
a	all of the above (an abbreviation for ugo)

u The user who owns the file (this means “you.”)

g The group the file belongs to.

o The other users

a all of the above (an abbreviation for ugo)

## Permissions

Of course, the permissions are the same letters that you see in the directory listing:

r	Permission to read the file.
w	Permission to write (or delete) the file.
x	Permission to execute the file, or, in the case of a directory, search it.

r Permission to read the file.

w Permission to write (or delete) the file.

x Permission to execute the file, or, in the case of a directory, search it.

Let’s say we have these files:

-rwxrwxrwx joe acctg wordmatic

-r--r--r-- joe acctg calcmatic

We’d like to remove write permission for the group and others on wordmatic, and add write and execute permission for all users on calcmatic.

Rather than try to figure out what the new permissions are and do these commands:

```
chmod go=rx wordmatic
```

```
chmod a=rwx calcmatic
```

The chmod command literally lets us add and subtract permissions from an existing set by using + or - instead of =.

Thus, we can take away the first file's write permission for the group and others with this command:

```
chmod go-w wordmatic
```

...and we can add write and execute permission to the second file for all users with:

```
chmod a+wx calcmatic
```

### Copying Permissions

As one other shortcut, it's possible to tell chmod "give users of one class the same permissions that some other class has."

Let's say we have these files:

```
d-----rwx joe acctg sales
```

```
-rw-r--r-- joe acctg info.dat
```

The other users have full permissions on the sales directory. We'd like to say "the user and group should be assigned (=) the permissions belonging to others." That translates to:

```
chmod ug=o
```

Similarly, to make info.dat readable and writable to the group, we can say:

```
chmod g=u info.dat
```

(you can read this as "the group is assigned (=) the permissions currently held by the user.")

You may also use + and - to add and subtract the permissions that currently belong to a different class of user.

You can't mix the standard permissions (r, w, and x) with the copying shortcuts. chmod will protest if you give it something like this:

```
chmod g=wu info.dat
```

### 4) umask = set file creation mask

The 'UMASK' is the default permission setting that is applied to your files and directories when they are created. After files and directories are created, the chmod command can be used to change the permissions to allow or disallow access as before. The UMASK is set when you login to a UNIX machine.

It is, however, possible to change your UMASK and put the UMASK in your login files so that your

default permissions are always set for files when you create them.

Just like chmod, a umask works on a number. However, instead of the numbers being ADDED like chmod, with a umask the numbers are SUBTRACTED from 7.

So from chmod –

Read – 4

Write – 2

Execute – 1

If a user wants all directories to be created with rwxr-xr-x, that is

Owner == Read, Write, Execute ==  $7 - 4 - 2 - 1 == 0$

Group == Read, Execute ==  $7 - 4 - 1 == 2$

Others == Read, Execute ==  $7 - 4 - 1 == 2$

Symbolic	Numeric	Description
-----	0000	no permissions
---x--x--x	0111	execute
--w--w--w-	0222	write
--wx-wx-wx	0333	write & execute
-r--r--r--	0444	read
-r-xr-xr-x	0555	read & execute
-rw-rw-rw-	0666	read & write
-rwxrwxrwx	0777	read, write, & execute
-rwxr-----	0740	user can read, write, & execute; group can only read; others have no permissions

## Chapter #3 - Test cases (task)

### Intro

Task verifies several candidate abilities including:

- ability to study independently by learning POSIX file systems standard
- perform data analysis by selecting important information
- design test cases and prepare documentation
- implement code based on design

### Test Task

Design a set of test cases for owner / permission / content modification testing of NFSv4 file system. Implement designed test cases as testing application (test suite). E.g, all tests are stored in “tests” folder and there is “main” file which run all tests and produces an output.

### The results of the task are:

1. Test documentation. Use the following format:

- Name of test case
- Description
- Steps
- Expected result of each step

2. Source code of test suite

3. Logs of the latest successful tests execution

### Test case example:

- Test name: Change file attributes to disable run, enable it, disable again.
- Description: test verifies that after several disable, enable actions permissions set to last value.

### Acceptance criteria:

1. At least 6 test cases have to be created
2. At least 2 test cases for ACL management verification (optional)

3. Test documentation

4. Use one of the following scripting language:

a. Python (preferable)    b. Ruby            c. Perl (in OOP style)

5. Test Suite has to prepare and clean environment

6. Keep logs in log file, Short summary should be printed at the end of testing. E.g.:

TC001: Passed    TC002: Failed            TC003: Passed

7. Should be executable at any Linux-Like system

8. Please use comments in the code



## Chapter #4 - Tests cases (Implementation)

### Goal:

To develop the test suite and create documentation (subject NFSv4 with ACL support test automatization for Linux-like systems [server-client sides])

### Test environment

Hostname	IP	Software	Description
fedora	192.168.100.182 (LAN)	Fedora 23 Workstation x86-x64 (ext4) + tools	Python 2.7, IDE Pycharm (create tests) Client NFSv4 (run tests)
rhel	192.168.100.176 (LAN)	Red Hat Enterprise Linux Server release 7.2 (Maipo) (RHEL) x86-x64 (LVM, XFS) + tools	Server NFSv4 (run tests)

### Daemons NFS

	both sides	server side
<b>user daemons</b>	<b>rpc.idmapd</b> - This process provides NFSv4 client and server upcalls which map between on-the-wire NFSv4 names (which are strings in the form of user@domain) and local UIDs and GIDs. For idmapd to function with NFSv4, the /etc/idmapd.conf must be configured. This service is required for use with NFSv4.	<b>rpc.nfsd</b> - Allows explicit NFS versions and protocols the server advertises to be defined. It works with the Linux kernel to meet the dynamic demands of NFS clients, such as providing server threads each time an NFS client connects. This process corresponds to the nfs service. <b>rpc.mountd</b> - daemon is still needed to handle the exports, but is not involved with network communication anymore (in other words, the client connects directly with the NFS daemon).
<b>kernel parts</b>	NFSv4, RPC, XDR, TCP, IPv4	-

### Packets and tools NFS

Name	Description	Addition
nfs-utils	The nfs-utils package provides a daemon for the kernel NFS server and related tools, which provides a much higher level of performance than the traditional Linux NFS server used by most users.	-
nfs4-acl-tools	The nfs4-acl-tools packages provide utilities for managing NFSv4 Access Control Lists (ACLs) on files and directories mounted on ACL-enabled NFSv4 file systems. These updated packages fix the following bug. This package contains commandline and GUI ACL utilities for the Linux NFSv4 client.	-
libnfsidmap	Is a library holding multiple methods of mapping names to id's and visa versa, mainly for NFSv4.	-
showmount	show mount information for an NFS server	-

## Configs NFS

File	Description	Comment
<b>Server side</b>		
/etc/exports	It is a main configuration file, controls which file systems are exported to remote hosts and specifies options. This file contains a list of entries; each entry indicates a volume that is shared and how it is shared.	There must be at least one entry with fsid=0. (this will be pseudo file system's /). <b>acl</b> option use.
/etc/sysconfig/nfs	This file is used to control which ports the required RPC services run on. Here the number of kernel threads, NFSv4 support and GSS security (kerberos) for NFS can be configured.	Not used on the project
/etc/idmapd.conf	It translates user and group ids into names, and to translate user and group names. Use to modify the default "Domain" to contain DNS domain name.	Not used on the project
<b>Client side</b>		
/etc/fstab	This file is used to control what file systems including NFS directories are mounted when the system boots.	<b>acl</b> option use
/etc/idmapd.conf	It translates user and group ids into names, and to translate user and group names. Use to modify the default "Domain" to contain DNS domain name.	Not used on the project

## Useful commands NFS

Description (Action)	Client side (Comand)	Server side (Comand)
Umount all nfs mounts on client	umount -a -t nfs	-
Check all nfs mounts on client	mount   grep nfs	-
Reexport of shares files	-	exportfs -r
Display of shares files	-	exportfs -v
Check all registered RPC programs (nfs, portmapper, mountd, ...)	rpcinfo -p	rpcinfo -p
Check mount information on NFS server	showmount -e <server IP>	showmount -e <server IP>
Display statistics kept about NFS client and server activity	nfsstat	nfsstat
Get file access control lists. For each file, getfacl displays the file name, owner, the group, and the Access Control List (ACL). If a directory has a default ACL, getfacl also displays the default ACL. Non-directories cannot have default ACLs.	getfacl -R <dir>  getfacl <file>	getfacl -R <dir>  getfacl <file>
This utility sets Access Control Lists (ACLs)	-	setfacl -R <options>

of files and directories.		
Display the NFSv4 Access Control List (ACL) for file (a file or directory), provided file is on a mounted NFSv4 filesystem which supports ACLs.	nfs4_getfacl <options>	nfs4_getfacl <options>
Manipulates the NFSv4 Access Control List (ACL) of one or more files (or directories), provided they are on a mounted NFSv4 filesystem which supports ACLs.	-	nfs4_setfacl <options>

## Install and settings of NFSv4 (\*\* without Kerberos)

### Server-side (hostname: rhel)

#### 1) Install the addition tools

The libraries and header files needed for Python developmen:

```
yum install python-devel
```

psutil - is a cross-platform library for retrieving information on running processes and system utilization (CPU, memory, disks, network) in Python:

```
pip install psutil
```

#### 2) Modify hosts file in order to resolve IP from hostname

```
vim /etc/hosts
```

```
192.168.100.176 rhel
```

```
192.168.100.182 fedora
```

```
ping fedora
```

```
PING fedora (192.168.100.182) 56(84) bytes of data.
```

```
64 bytes from fedora (192.168.100.182): icmp_seq=1 ttl=64 time=0.335 ms
```

#### 3) In order to use ACLs enable mount option

```
vim /etc/fstab
```

```
dev/mapper/rhel_nfs-root/          xfs  defaults,acl    0 0
```

```
mount -a
```

#### 4) Check which loadable kernel modules NFS currently loaded

```
lsmod | grep nfs
```

```
nfs          251815  0
```

```
fscache      64987  1 nfs
```

```
nfsd         302351  1
```

```
auth_rpcgss  59314  1 nfsd
```

```
nfs_acl      12837  1 nfsd
```

```
lockd        93572  2 nfs,nfsd
```

```
grace        13288  2 nfsd,lockd
```

```
sunrpc       300421  8 nfs,nfsd,auth_rpcgss,lockd,nfs_acl
```

#### 5) Check and show information about NFS modules:

```
modinfo nfs
```

```
filename:      /lib/modules/3.10.0-327.13.1.el7.x86_64/kernel/fs/nfs/nfs.ko
```

```
modinfo nfsv3
```

```
filename: /lib/modules/3.10.0-327.13.1.el7.x86_64/kernel/fs/nfs/nfsv3.ko
```

```
modinfo nfsv4
```

```
filename: /lib/modules/3.10.0-327.13.1.el7.x86_64/kernel/fs/nfs/nfsv4.ko
```

```
modinfo nfsd
```

```
filename: /lib/modules/3.10.0-327.13.1.el7.x86_64/kernel/fs/nfsd/nfsd.ko
```

```
modinfo nfs_acl
```

```
filename: /lib/modules/3.10.0-327.13.1.el7.x86_64/kernel/fs/nfs_common/nfs_acl.ko
```

```
modinfo nfs_layout_flexfiles
```

```
filename: /lib/modules/3.10.0-327.13.1.el7.x86_64/kernel/fs/nfs/flexfilelayout/nfs_layout_flexfiles.ko
```

```
modinfo nfs_layout_nfsv41_files
```

```
filename: /lib/modules/3.10.0-327.13.1.el7.x86_64/kernel/fs/nfs/filelayout/nfs_layout_nfsv41_files.ko
```

## 6) Check linux kernel for ACL support (find **\*=Y** option in accordance with the task)

```
uname -a
```

```
Linux rhel 3.10.0-327.13.1.el7.x86_64 #1 SMP Mon Feb 29 13:22:02 EST 2016 x86_64 x86_64 x86_64 GNU/Linux
```

```
grep -i acl /boot/config*
```

```
/boot/config-3.10.0-327.13.1.el7.x86_64:CONFIG_EXT4_FS_POSIX_ACL=y
```

```
/boot/config-3.10.0-327.13.1.el7.x86_64:CONFIG_XFS_POSIX_ACL=y
```

```
/boot/config-3.10.0-327.13.1.el7.x86_64:CONFIG_BTRFS_FS_POSIX_ACL=y
```

```
/boot/config-3.10.0-327.13.1.el7.x86_64:CONFIG_FS_POSIX_ACL=y
```

```
/boot/config-3.10.0-327.13.1.el7.x86_64:CONFIG_GENERIC_ACL=y
```

```
/boot/config-3.10.0-327.13.1.el7.x86_64:CONFIG_TMPFS_POSIX_ACL=y
```

```
/boot/config-3.10.0-327.13.1.el7.x86_64:CONFIG_NFS_V3_ACL=y
```

```
/boot/config-3.10.0-327.13.1.el7.x86_64:CONFIG_NFSD_V2_ACL=y
```

```
/boot/config-3.10.0-327.13.1.el7.x86_64:CONFIG_NFSD_V3_ACL=y
```

```
/boot/config-3.10.0-327.13.1.el7.x86_64:CONFIG_NFS_ACL_SUPPORT=m
```

```
/boot/config-3.10.0-327.13.1.el7.x86_64:CONFIG_CIFS_ACL=y
```

If there is N instead of Y, then it means linux kernel doesn't support ACL and need to be recompiled in accordance with the task.

## 7) Install NFS server and tools

```
yum install nfs-utils nfs4-acl-tools libnfsidmap
```

## 8) Check and enable NFS server services

```
systemctl list-unit-files | grep nfs
```

```
proc-fs-nfsd.mount          static
```

```
var-lib-nfs-rpc_pipefs.mount static
```

<i>nfs-blkmap.service</i>	<i>disabled</i>
<i>nfs-config.service</i>	<i>static</i>
<i>nfs-idmap.service</i>	<i>static</i>
<i>nfs-idmapd.service</i>	<i>static</i>
<i>nfs-lock.service</i>	<i>static</i>
<i>nfs-mountd.service</i>	<i>static</i>
<i>nfs-rquotad.service</i>	<i>disabled</i>
<i>nfs-secure-server.service</i>	<i>static</i>
<i>nfs-secure.service</i>	<i>static</i>
<b><i>nfs-server.service</i></b>	<b><i>disabled</i></b>
<i>nfs-utils.service</i>	<i>static</i>
<b><i>nfs.service</i></b>	<b><i>disabled</i></b>
<i>nfslock.service</i>	<i>static</i>
<i>nfs-client.target</i>	<i>enabled</i>

*systemctl enable nfs-server.service*

*systemctl enable nfs.service*

*systemctl start nfs-server.service*

*systemctl start nfs.service*

**9) Create NFS share and change permissions (the export filesystem) - a directories to share with client servers**

*mkdir -p /nfs*

*chmod a+rwxt /export*

**10) Share directories of NFS server for any (LAN, WAN, ...). Exports - NFS server export table.**

*vim /etc/exports*

*/nfs \*(rw,fsid=0,nohide, no\_root\_squash, insecure,no\_subtree\_check,sync)*

**/nfs**

**\*** - users from any IP address of client machine are allowed to mount directories (means any client)

**rw** - allow both read and write requests on this NFS volume. The default is to disallow any request which changes the filesystem.

**fsid=0** - export a directory over NFSv4. NFSv4 has a concept of a root of the overall exported filesystem. The export point exported with fsid=0 will be used as this root. The /nfs directory will be root for clients. For example, if you got /nfs/tests subdirectory, then client would see them as /tests directory. NFS needs to be able to identify each filesystem that it exports. For NFSv4, there is a distinguished filesystem which is the root of all exported filesystem. This is specified with fsid=root or fsid=0 both of which mean exactly the same thing. Only for “root” directory.

**nohide** - setting the nohide option on a filesystem causes it not to be hidden, and an appropriately

authorised client will be able to move from the parent to that filesystem without noticing the change.

**no\_root\_squash** - turn off root squashing. This option is mainly useful for diskless clients. By default, any file request made by user root on the client machine is treated as by user nobody on the server. (Exactly which UID the request is mapped to depends on the UID of user “nobody” on the server, not the client.) If no\_root\_squash is selected, then root on the client machine will have the same level of access to the files on the system as root on the server.

**insecure** - option in this entry allows clients with NFS implementations that don't use a reserved port for NFS.

**no\_subtree\_check** - this option disables subtree checking, which has mild security implications, but can improve reliability in some circumstances. If a subdirectory of a filesystem is exported, but the whole filesystem isn't then whenever a NFS request arrives, the server must check not only that the accessed file is in the appropriate filesystem (which is easy) but also that it is in the exported tree (which is harder). This check is called the subtree\_check.

**sync** - reply to requests only after the changes have been committed to stable storage (if async - improve performance, but at the cost that an unclean server restart (i.e. a crash) can cause data to be lost or corrupted). All changes to the according filesystem are immediately flushed to disk; the respective write operations are being waited for.

#### 11) Restart NFS server services

```
systemctl restart nfs-server.service
```

```
systemctl restart nfs.service
```

#### 12) Reexport all directories after modifying **/etc/exports** and display a list of shares files and export options on a NFS server

```
exportfs -r
```

```
exportfs -v
```

```
/nfs<world>(rw,wdelay,nohide,insecure,no_root_squash,no_subtree_check,fsid=0,sec=sys,rw,insecure,no_root_squash,no_all_squash)
```

```
/nfs/tests<world>(rw,wdelay,nohide,insecure,root_squash,no_subtree_check,fsid=0,sec=sys,rw,insecure,root_squash,no_all_squash)
```

#### 13) Config or Disable firewall (firewalld or iptables services) on NFS server to allow client servers to access NFS shares. Open TCP port # 2049 which is used by NFSv4.

##### a) Config firewall

```
firewall-cmd --permanent --add-service nfs *** need only one for remote mount via TCP port 2049
```

```
firewall-cmd --permanent --add-service rpc-bind
```

```
firewall-cmd --permanent --add-service mountd
```

```
firewall-cmd --reload
```

```
firewall-cmd --list-all
```

```
public (default, active)
```

```
interfaces: ens192
```

*sources:*

*services: dhcpv6-client mountd nfs rpc-bind ssh*

*ports:*

*masquerade: no*

*forward-ports:*

*icmp-blocks:*

*rich rules:*

*cat /etc/services | grep mountd*

*mountd 20048/tcp # NFS mount protocol*

*mountd 20048/udp # NFS mount protocol*

*cat /etc/services | grep nfs*

*nfs 2049/tcp nfsd shilp # Network File System*

*nfs 2049/udp nfsd shilp # Network File System*

*nfs 2049/sctp nfsd shilp # Network File System*

*cat /etc/services | grep rpcbind*

*sunrpc 111/tcp portmapper rpcbind # RPC 4.0 portmapper TCP*

*sunrpc 111/udp portmapper rpcbind # RPC 4.0 portmapper UDP*

## **b) Disable firewall**

*systemctl disable firewalld*

*systemctl stop firewalld*

*systemctl status firewalld*

*systemctl status firewalld*

● *firewalld.service - firewalld - dynamic firewall daemon*

*Loaded: loaded (/usr/lib/systemd/system/firewalld.service; disabled; vendor preset: enabled)*

*Active: inactive (dead)*

## **14) Disable and SELinux**

*vim /etc/selinux/config*

*SELINUX=disable*

*\*\*\**

*sestatus*

*SELinux status: disabled*

## **15) Check all registered RPC programs (nfs, portmapper, mountd)**



*rpcinfo -p*

```
program vers proto  port  service
100000  4  tcp   111  portmapper
100000  4  udp   111  portmapper
100024  1  udp  33569  status
100024  1  tcp  41527  status
100005  1  udp  20048  mountd
100005  1  tcp  20048  mountd
100003  4  tcp   2049  nfs
100003  4  udp   2049  nfs
100227  3  tcp   2049  nfs_acl
100227  3  udp   2049  nfs_acl
100021  1  udp  39271  nlockmgr
100021  1  tcp  60402  nlockmgr
```

#### **16) Set UNIX extended ACLs (add r/w permissions for user: he on directory: /nfs)**

\*\*\*By default, if the file system being exported by an NFSv4 server supports ACLs and the NFS client can read ACLs, ACLs are utilized by the client system.

*setfacl -R -m u:he:rwX /nfs*

#### **17) Check UNIX extended ACLs**

*getfacl -R /nfs*

\*\*\*By default, if the file system being exported by an NFSv4 server supports ACLs and the NFS client can read ACLs, ACLs are utilized by the client system.

*getfacl: Removing leading '/' from absolute path names*

**# file: nfs**

*# owner: root*

*# group: root*

*user::rwX*

***user:he:rwX***

*group::r-x*

*mask::rwx*

*other::r-x*

***# file: nfs/tests***

*# owner: root*

*# group: root*

*user::rwx*

***user:he:rwx***

*group::r-x*

*mask::rwx*

*other::r-x*

***# file: nfs/tests/music\_for\_programming\_00-manifesto.mp3***

*# owner: he*

*# group: he*

*user::rwx*

***user:he:rwx***

*group::rwx*

*mask::rwx*

*other::rwx*

***# file: nfs/DDNTestTaskE1-E2-true.pdf***

*# owner: he*

*# group: he*

*user::rwx*

***user:he:rwx***

*group::rwx*

*mask::rwx*

*other::r-x*

**18) Install and settings addition software need for run testcases**

**a) rsh - remote shell access (need in order to receive commands from remote client)**

## Install rsh and rshd:

```
yum install rsh rsh-server
```

```
rpm -qa | grep rsh
```

```
rsh-server-0.17-76.el7_1.1.x86_64
```

```
rsh-0.17-76.el7_1.1.x86_64
```

## Start rsh-server daemons:

```
systemctl enable rsh.socket
```

```
systemctl enable rlogin.socket
```

```
systemctl enable rexec.socket
```

```
systemctl start rsh.socket
```

```
systemctl start rlogin.socket
```

```
systemctl start rexec.socket
```

```
systemctl status rsh.socket
```

- *rsh.socket - Remote Shell Facilities Activation Socket*

*Loaded: loaded (/usr/lib/systemd/system/rsh.socket; enabled; vendor preset: disabled)*

*Active: active (listening)*

```
systemctl status rlogin.socket
```

- *rlogin.socket - Remote Login Facilities Activation Socket*

*Loaded: loaded (/usr/lib/systemd/system/rlogin.socket; enabled; vendor preset: disabled)*

*Active: active (listening)*

```
systemctl status rexec.socket
```

- *rexec.socket - Remote Execution Facilities Activation Socket*

*Loaded: loaded (/usr/lib/systemd/system/rexec.socket; enabled; vendor preset: disabled)*

*Active: active (listening)*

## Configure rsh-server:

*vim /root/.rhosts* - allow the user root on the client fedora to log in as root on the target (server)

*fedora root*

*vim /etc/securetty* - enable external root user to execute the command (lists terminals from which root can log in)

*rsh*

*rexec*

*rlogin*

## Client-side (hostname: fedora)

### 1) Install the addition tools

The libraries and header files needed for Python developmen:

```
yum install python-devel
```

psutil - is a cross-platform library for retrieving information on running processes and system utilization (CPU, memory, disks, network) in Python:

```
pip install psutil
```

### 2) Modify hosts file in order to resolve IP from hostname

```
vim /etc/hosts
```

```
192.168.100.176 rhel
```

```
192.168.100.182 fedora
```

```
ping rhel
```

```
PING rhel (192.168.100.176) 56(84) bytes of data.
```

```
64 bytes from rhel (192.168.100.176): icmp_seq=1 ttl=64 time=0.594 ms
```

```
***
```

### 3) Disable firewall

```
systemctl disable firewalld
```

```
systemctl stop firewalld
```

```
systemctl status firewalld
```

```
systemctl status firewalld
```

```
● firewalld.service - firewalld - dynamic firewall daemon
```

```
Loaded: loaded (/usr/lib/systemd/system/firewalld.service; disabled; vendor preset: enabled)
```

```
Active: inactive (dead)
```

### 4) Disable and SELinux

```
vim /etc/selinux/config
```

```
SELINUX=disable
```

```
***
```

```
sestatus
```

```
SELinux status: disabled
```

### 5) Install nfs utils for client NFS service

```
yum install nfs-utils nfs4-acl-tools libnfsidmap
```

## 6) Enable and start NFS service

```
systemctl enable nfs-client.target
```

```
systemctl start nfs-client.target
```

## 7) Check nfs shares to the clients on NFS server

```
showmount -e 192.168.100.176
```

Export list for 192.168.100.176:

```
/nfs/tests *
```

```
/nfs *
```

## 8) Mount the exported file system

```
mount -t nfs4 192.168.100.176:/ /nfs
```

Observe that only "/" is given instead of the actual exported path name

## 9) Check mounted NFSv4 file system

```
mount | grep nfs
```

```
nfsd on /proc/fs/nfsd type nfsd (rw,relatime)
```

```
sunrpc on /var/lib/nfs/rpc_pipefs type rpc_pipefs (rw,relatime)
```

```
192.168.100.176:/ on /nfs type nfs4
```

```
(rw,relatime,vers=4.1,rsize=524288,wsiz=524288,namlen=255,hard,proto=tcp,timeo=600,retrans=2,sec=sys,clientaddr=192.168.100.182,local_lock=none,addr=192.168.100.176) *** ver 4.1
```

```
df -hT | grep nfs
```

```
192.168.100.176:/ nfs4 37G 5.5G 32G 15% /nfs
```

## 10) Mount NFS file system permanently in order to mount after reboot system

```
vim /etc/fstab
```

```
192.168.100.176:/ /nfs nfs4_netdev,auto 0 0
```

```
mount -a
```

## 11) Check all registered RPC programs

```
sdrpcinfo -p
```

```
program vers proto port service
```

```
100000 4 tcp 111 portmapper
```

```
100000 4 udp 111 portmapper
```

## 12) Install and settings additional software needed for running testcases

a) rsh - remote shell access (needed in order to execute commands on remote server)

**Install rsh and rshd:**

```
yum install rsh rsh-server
```

```
rpm -qa | grep rsh
```

```
rsh-server-0.17-76.el7_1.1.x86_64
```

```
rsh-0.17-76.el7_1.1.x86_64
```

**13) Check UNIX extended ACLs**

```
getfacl -R /nfs
```

\*\*\*By default, if the file system being exported by an NFSv4 server supports UNIX extended ACLs and the NFS client can read ACLs, ACLs are utilized by the client system.

Test #1: NFSv4 - Test the maximum number of ACEs (Access Control Entries) supported by file system [Extended ACLs for UNIX]

Limits for the max number of the ACEs per individual file/directory:

- EXT2, EXT3, EXT4: 32
- XFS: 25
- GFS2: 25

Auto check file system type for NFSv4 server export directory and check limits for max count of ACEs [Extended ACLs for UNIX]

Test #2: NFSv4 - Stress test for a large number of random operations setting ACEs [Extended ACLs for UNIX]

The testing of correct operation of the services of NFSv4 server with a large number of random operations setting ACEs (Access Control Entries)

Part 1 - Emulate a large number random operations with ACEs for export directory on the NFSv4 server, health check NFSv4 service

Part 2 - Emulate a large number random operations with ACEs for random files in export directory on the NFSv4 server, health check NFSv4 service

Test #3: NFSv4 - Stress test for a large number operations:

[NFSv4 server side] : export / unexport  
[NFSv4 client side] : mount / unmount

Part 1 - Emulation of export a large number of directories on the NFSv4 server and their subsequent mounting on the NFSv4 client, health check NFSv4 service

Part 2 - Emulation of a large number of random operations with exports on NFSv4 server and mounts on NFSv4 client, health check NFSv4 service

\*\*\* The test data will be created in the directory /mnt/\* from both sides NFSv4 client and server

Test #4: The complex test for owner/permission/content modification of NFSv4 file system [POSIX]

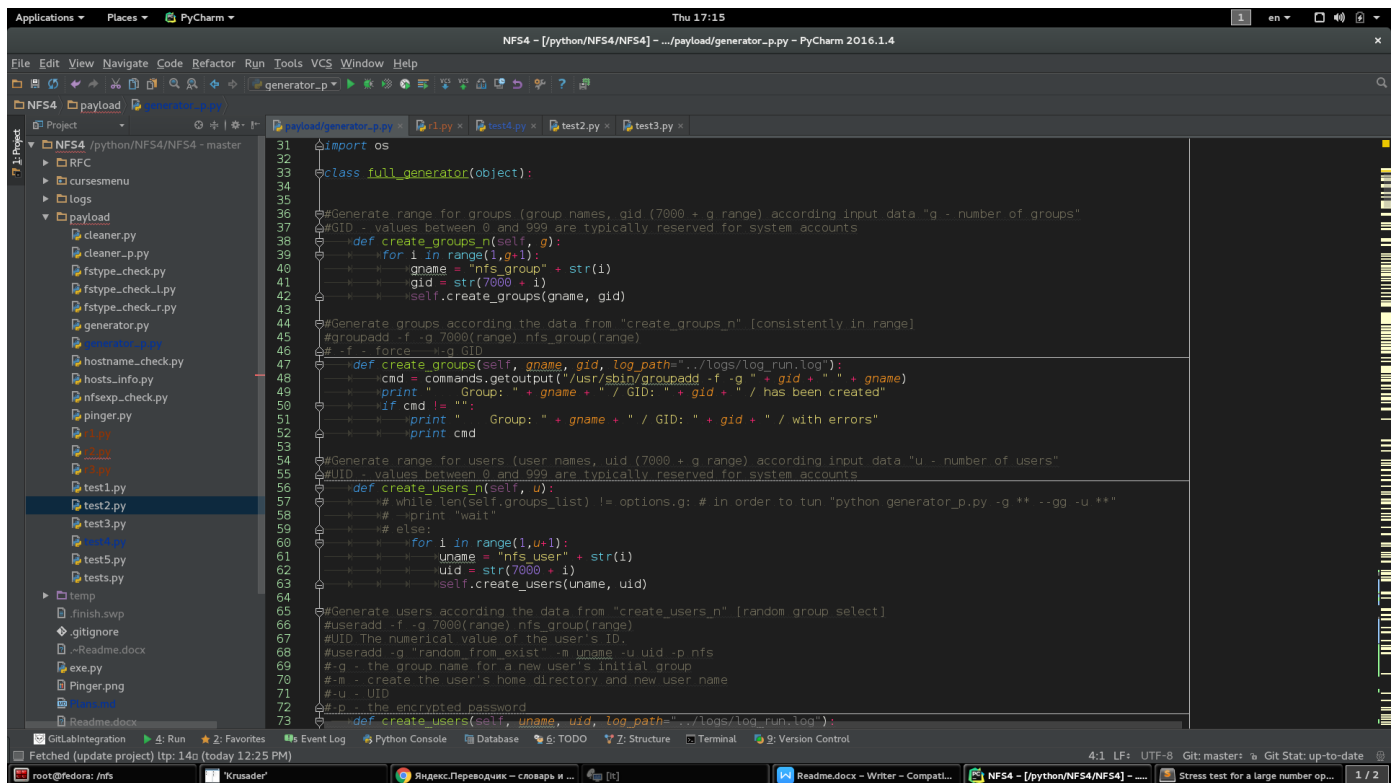
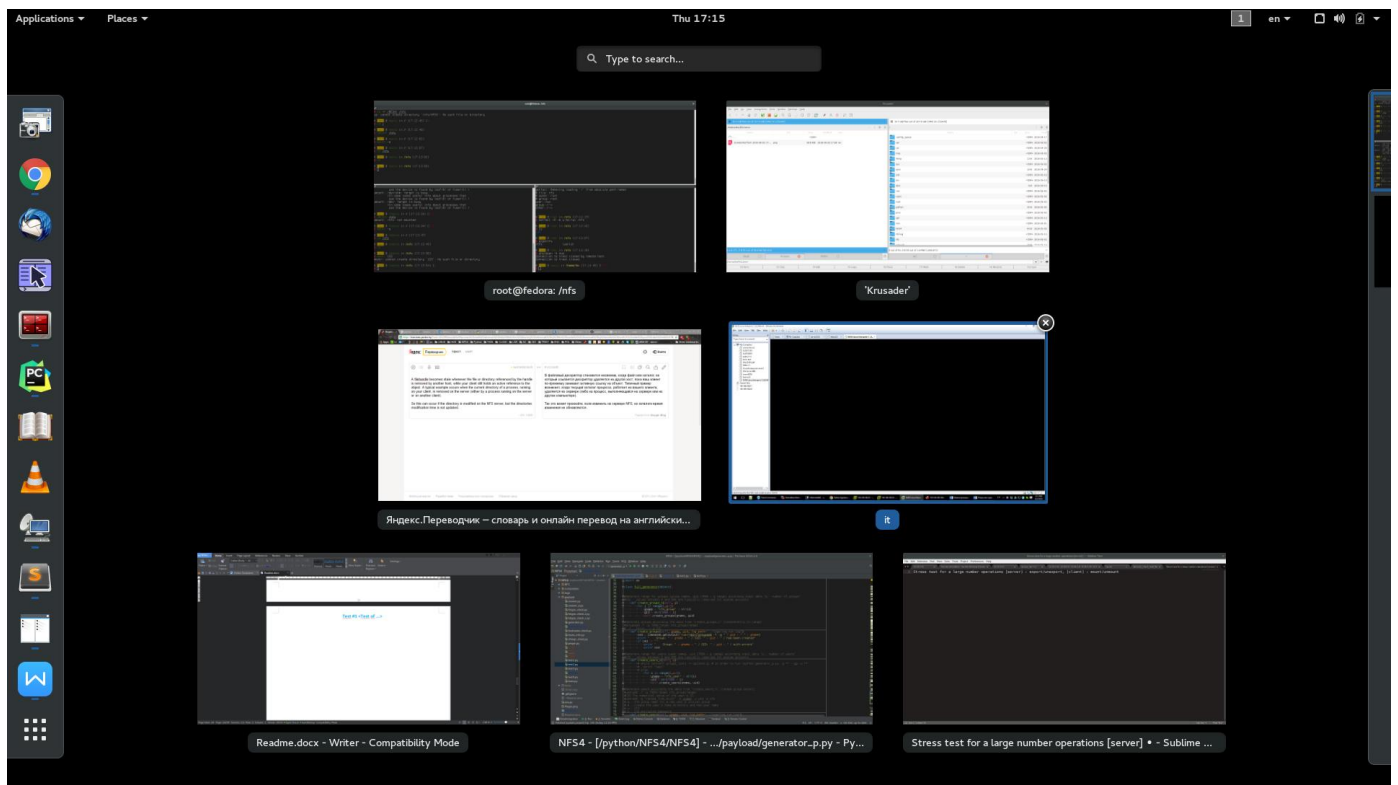
Part 1 - Check the ability of check the permissions of the files

> Description : Check the ability of changes the permissions of the files  
> Steps : Run used functions of full\_generator  
> Expected result : Display "TEST #4 PART 2 HAS BEEN PASSED"

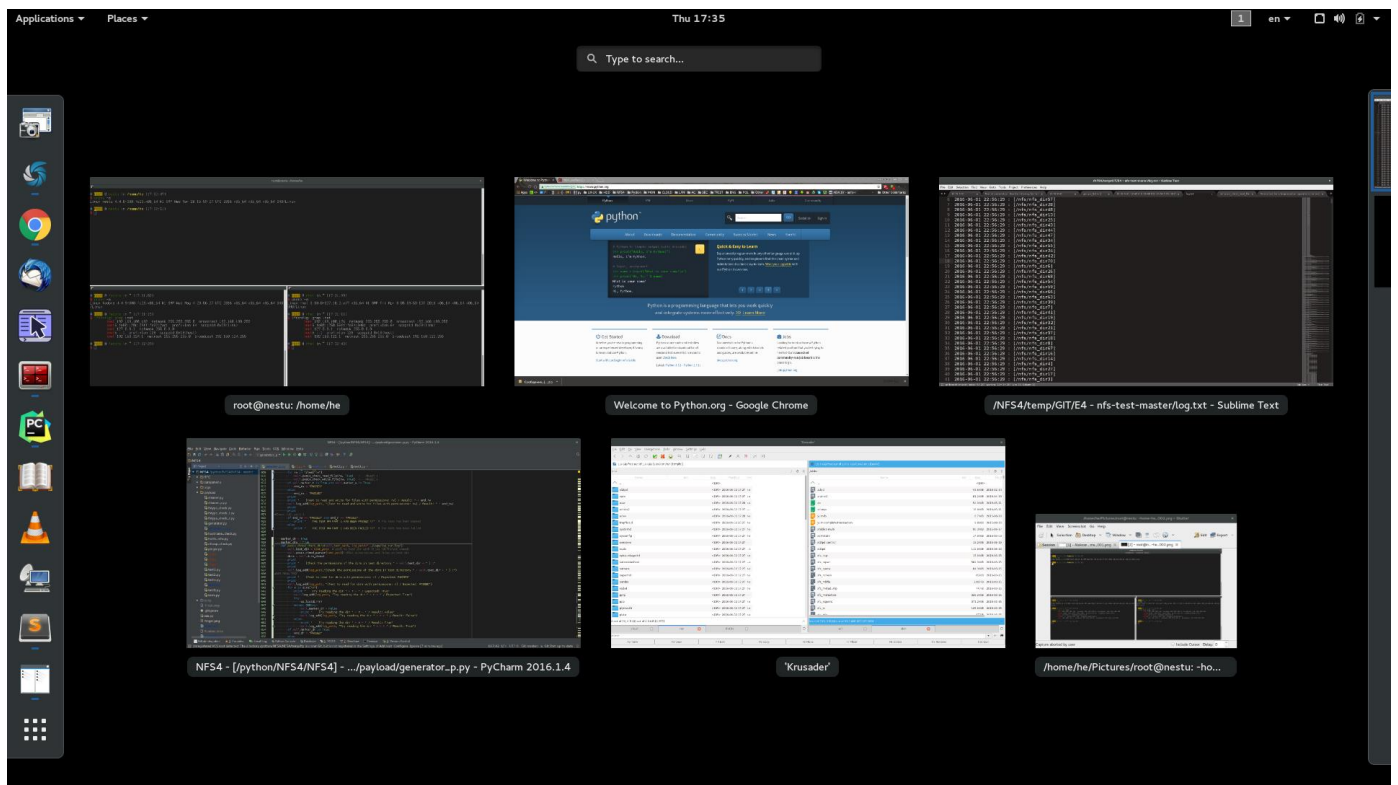
Part 2 - Check the ability of check the permissions of the directories

> Description : Check the ability of changes the permissions of the directories  
> Steps : Run used functions of full\_generator  
> Expected result : Display "TEST #4 PART 2 HAS BEEN PASSED"

\*\*\* All parts of the complex test will be performed sequentially and will be running in the export directory on NFSv4 server side







```
<NFSv4 test suite> [Run from Client side]
```

#### MENU:

- 1 - Read Help
- 2 - Check test environment
- 3 - Run tests
- 4 - Exit

```
<NFSv4 test suite> [Run from Client side]
```

#### Run tests:

- 1 - 1. Test the maximum number of ACEs supported by file system [Extended ACLs for UNIX]
- 2 - 2. Stress test for a large number of random operations setting ACEs [Extended ACLs for UNIX]
- 3 - 3. Stress test for a large number of operations [server] : export/unexport, [client] : mount/unmount
- 4 - 4. Complex test for owner/permission/content modification of NFSv4 file system
- 5 - 5. Stress test for a large number of random operations setting ACEs [NFSv4 ACLs]
- 6 - Run all tests 1..5
- 7 - Watch the log
- 8 - Return to <NFSv4 test suite> [Run from Client side] menu

```
! [main window of program](https://github.com/AleksNestu/NFSv4/blob/master/Screen.png)
```

#### LICENSE

```
=====
Free for personal use:)
```

#### 1. The brief overview of the project:

```
=====
Test cases for automatic testing of NFSv4 file system on Linux-like systems.
Subject of testing: content, ACLs (Access Control Lists), etc.
```

#### 2. Project description:

```
=====
The structure of the project:
- "./RFC/" - the directory which contains RFC documents about NFSv4 and ACLs;
- "./cursesmenu/" - the directory which contains python module "The simple console menu system using the curses library";
- "./logs/" - the directory which contains the execution results (log-files in text format) [log_run.log, log_result.log];
- "./payload/" - the directory which contains a set of tests (each file is a separate test) *** can be supplemented;
- "./Readme.pdf" - the full description of my work with the project and useful info [Intro, NFSv4, ACLs, Test task, Test Implementation];
- "./Screen.png" - the program's appearance;
```

## Test #1: NFSv4 - Test the maximum number of ACEs (Access Control Entries) supported by file system [Extended ACLs for UNIX]

=====

Test #1: NFSv4 - Test the maximum number of ACEs (Access Control Entries)  
supported by file system [Extended ACLs for UNIX]

Limits for the max number of the ACEs per individual file/directory:

- EXT2, EXT3, EXT4: 32
- XFS: 25
- GFS2: 25

Auto check file system type for NFSv4 server export directory and check  
limits for max count of ACEs [Extended ACLs for UNIX]

=====

In order to run the test enter the required data:

1) Hostname of the NFSv4 server:

[input] : rhel  
hostname have resolved  
NFSv4 server: rhel

2) Path to the exported directory on NFSv4 server (/nfsdir, /mnt/nfs, ...):

[input] : /nfs  
NFS server's exported directory is existing  
NFSv4 exported dir: /nfs

3) NFSv4 server exported dir filesystem type & ACEs limits:

NFSv4 server fs type for exp dir /nfs : xfs

ACEs max count (UNIX extended ACLs): 25

4) The number of users and groups to be created on the NFSv4 server:

The number of users [25..999] [input] : 30

The number of groups [25..999] [input] : 25

5) The directory and file which will be created on the export directory on the NFSv4 server:

Test directory (folder, dir, ...) in NFSv4 server export dir [input] : dir

Test file (file, goal, ...) in NFSv4 server export dir [input] : file

Great! The input data have been received!

After 5 seconds the test will be started...

=====

Client : ('Linux 4.4.9-300.fc23.x86\_64 x86\_64\n', None)

Hostname : fedora

IP : 192.168.100.182

Server : ('Linux 3.10.0-327.18.2.el7.x86\_64 x86\_64\n', None)

Hostname : rhel

IP : 192.168.100.176

Ping : fedora --->>> rhel

PING rhel (192.168.100.176) 56(84) bytes of data.

64 bytes from rhel (192.168.100.176): icmp\_seq=1 ttl=64 time=0.300 ms

64 bytes from rhel (192.168.100.176): icmp\_seq=2 ttl=64 time=0.340 ms

64 bytes from rhel (192.168.100.176): icmp\_seq=3 ttl=64 time=0.334 ms

--- rhel ping statistics ---

3 packets transmitted, 3 received, 0% packet loss, time 1999ms

rtt min/avg/max/mdev = 0.300/0.324/0.340/0.027 ms

[Create 25 groups on the NFSv4 server] :

Group: nfs\_group1 / GID: 7001 / has been created

Group: nfs\_group2 / GID: 7002 / has been created

...

Group: nfs\_group25 / GID: 7025 / has been created

[Create 30 users on the NFSv4 server] :

User: nfs\_user1 / UID: 7001 / GID: nfs\_group20 / has been created

User: nfs\_user2 / UID: 7002 / GID: nfs\_group12 / has been created

...

User: nfs\_user28 / UID: 7028 / GID: nfs\_group18 / has been created

User: nfs\_user29 / UID: 7029 / GID: nfs\_group6 / has been created

User: nfs\_user30 / UID: 7030 / GID: nfs\_group24 / has been created

[Create the directory and file on the NFSv4 server export directory] :

Test directory on the [rhel] : /nfs/dir has been created

Test file on the [rhel] : /nfs/file has been created

[Test the maximum number of ACEs for directory] :

ACE # 5 : setfacl -m u:nfs\_user24:x /nfs/dir

ACE # 6 : setfacl -m u:nfs\_user26:r /nfs/dir

ACE # 7 : setfacl -m u:nfs\_user21:x /nfs/dir

ACE # 8 : setfacl -m u:nfs\_user4:xr /nfs/dir

ACE # 9 : setfacl -m u:nfs\_user28:w /nfs/dir

ACE # 10 : setfacl -m u:nfs\_user18:xr /nfs/dir

ACE # 11 : setfacl -m u:nfs\_user12:xr /nfs/dir

ACE # 12 : setfacl -m u:nfs\_user20:r /nfs/dir

ACE # 13 : setfacl -m u:nfs\_user3:xr /nfs/dir

ACE # 14 : setfacl -m u:nfs\_user2:x /nfs/dir

ACE # 15 : setfacl -m u:nfs\_user9:w /nfs/dir

ACE # 16 : setfacl -m u:nfs\_user8:xr /nfs/dir

ACE # 17 : setfacl -m u:nfs\_user11:wr /nfs/dir

ACE # 18 : setfacl -m u:nfs\_user10:xr /nfs/dir

ACE # 19 : setfacl -m u:nfs\_user27:wr /nfs/dir

ACE # 20 : setfacl -m u:nfs\_user16:wr /nfs/dir

ACE # 21 : setfacl -m u:nfs\_user1:xr /nfs/dir

ACE # 22 : setfacl -m u:nfs\_user7:xr /nfs/dir  
ACE # 23 : setfacl -m u:nfs\_user14:xwr /nfs/dir  
ACE # 24 : setfacl -m u:nfs\_user15:x /nfs/dir  
ACE # 25 : setfacl -m u:nfs\_user30:xr /nfs/dir  
ACE # 26 : setfacl -m u:nfs\_user5:wr /nfs/dir - ERROR!!!

setfacl: /nfs/dir: Argument list too long

Reached the maximum number of ACEs: 25

THE TEST #1 HAS BEEN PASSED!!!

[Test the maximum number of ACEs for file] :

ACE # 5 : setfacl -m u:nfs\_user3:r /nfs/file  
ACE # 6 : setfacl -m u:nfs\_user21:wr /nfs/file  
ACE # 7 : setfacl -m u:nfs\_user28:xr /nfs/file  
ACE # 8 : setfacl -m u:nfs\_user29:r /nfs/file  
ACE # 9 : setfacl -m u:nfs\_user27:wr /nfs/file  
ACE # 10 : setfacl -m u:nfs\_user15:wr /nfs/file  
ACE # 11 : setfacl -m u:nfs\_user20:w /nfs/file  
ACE # 12 : setfacl -m u:nfs\_user12:x /nfs/file  
ACE # 13 : setfacl -m u:nfs\_user25:xwr /nfs/file  
ACE # 14 : setfacl -m u:nfs\_user7:r /nfs/file  
ACE # 15 : setfacl -m u:nfs\_user13:xwr /nfs/file  
ACE # 16 : setfacl -m u:nfs\_user30:wr /nfs/file  
ACE # 17 : setfacl -m u:nfs\_user9:xr /nfs/file  
ACE # 18 : setfacl -m u:nfs\_user18:w /nfs/file  
ACE # 19 : setfacl -m u:nfs\_user2:xwr /nfs/file  
ACE # 20 : setfacl -m u:nfs\_user16:xwr /nfs/file  
ACE # 21 : setfacl -m u:nfs\_user10:wr /nfs/file  
ACE # 22 : setfacl -m u:nfs\_user1:wr /nfs/file  
ACE # 23 : setfacl -m u:nfs\_user23:wr /nfs/file  
ACE # 24 : setfacl -m u:nfs\_user19:wr /nfs/file  
ACE # 25 : setfacl -m u:nfs\_user8:xr /nfs/file  
ACE # 26 : setfacl -m u:nfs\_user5:x /nfs/file - ERROR!!!

setfacl: /nfs/file: Argument list too long

Reached the maximum number of ACEs: 25

THE TEST #1 HAS BEEN PASSED!!!

[Clean created directory and file on the NFSv4 server] :

Test directory on the [rhel] : /nfs/dir has been removed

Test file on the [rhel] : /nfs/file has been removed

[Clean created users and groups on the NFSv4 server] :

User del: nfs\_user1 / has been done

...

User del: nfs\_user30 / has been done

Group del: nfs\_group1 / has been done

...

Group del: nfs\_group24 / has been done

Group del: nfs\_group25 / has been done

#####

Test #1 NFSv4 maximum number of ACEs [Extended ACLs for UNIX]

[PASSED] [PASSED] [PASSED] [PASSED] [PASSED] [PASSED] [PASSED]

In order to get more information:

./logs/log\_run.log" - execution log (detailed information)

./logs/log\_result.log - log with the results

#####

## Test #2 NFSv4 - Stress test for a large number of random operations setting ACEs [Extended ACLs for UNIX] [PASSED]

=====

**Test #2: NFSv4 - Stress test for a large number of random operations  
setting ACEs [Extended ACLs for UNIX]**

**The testing of correct operation of the services of NFSv4 server  
with a large number of random operations setting ACEs  
(Access Control Entries)**

**Part 1 - Emulate a large number random operations with ACEs for export  
directory on the NFSv4 server, health check NFSv4 service**

**Part 2 - Emulate a large number random operations with ACEs for random  
files in export directory on the NFSv4 server, health check  
NFSv4 service**

=====

**In order to run the test enter the required data:**

**1) Hostname of the NFSv4 server:**

**[input] : rhel**  
**hostname have resolved**  
**NFSv4 server: rhel**

**2) Path to the exported directory on NFSv4 server (/nfsdir, /mnt/nfs, ...):**

**[input] : /nfs**  
**NFS server's exported directory is existing**  
**NFSv4 exported dir: /nfs**

**3) The number of users and groups to be created on the NFSv4 server:**

The number of users [50..1000] [input] : 50

The number of groups [50..1000] [input] : 60

4) The number of files to be created in the export dir on the NFSv4 server:

The number of files [10..50] [input] : 20

5) The number of cycles to perform random operation in the test:

The number of cycles [500..10000] [input] : 500

Great! The input data have been received!

After 5 seconds the test will be started...

=====

Client : ('Linux 4.4.9-300.fc23.x86\_64 x86\_64\n', None)

Hostname : fedora

IP : 192.168.100.182

Server : ('Linux 3.10.0-327.18.2.el7.x86\_64 x86\_64\n', None)

Hostname : rhel

IP : 192.168.100.176

Ping : fedora --->>> rhel

PING rhel (192.168.100.176) 56(84) bytes of data.

64 bytes from rhel (192.168.100.176): icmp\_seq=1 ttl=64 time=0.276 ms

64 bytes from rhel (192.168.100.176): icmp\_seq=2 ttl=64 time=0.361 ms

64 bytes from rhel (192.168.100.176): icmp\_seq=3 ttl=64 time=0.365 ms

--- rhel ping statistics ---

3 packets transmitted, 3 received, 0% packet loss, time 2000ms

rtt min/avg/max/mdev = 0.276/0.334/0.365/0.041 ms

[Create 60 groups on the NFSv4 server] :

Group: nfs\_group1 / GID: 7001 / has been created



...

Group: nfs\_group60 / GID: 7060 / has been created

[Create 50 users on the NFSv4 server] :

User: nfs\_user1 / UID: 7001 / GID: nfs\_group56 / has been created

...

User: nfs\_user50 / UID: 7050 / GID: nfs\_group11 / has been created

[Part 1 - Stress test for a large number random ACEs for export directory on the NFSv4 server] :

ACE # 1 : setfacl -m u:nfs\_user21:wr /nfs & setfacl -x u:nfs\_user21 /nfs

...

ACE # 497 : setfacl -m u:nfs\_user8:xwr /nfs & setfacl -x u:nfs\_user8 /nfs

ACE # 498 : setfacl -m u:nfs\_user44:xw /nfs & setfacl -x u:nfs\_user44 /nfs

ACE # 499 : setfacl -m u:nfs\_user28:r /nfs & setfacl -x u:nfs\_user28 /nfs

ACE # 500 : setfacl -m u:nfs\_user11:xr /nfs & setfacl -x u:nfs\_user11 /nfs

THE TEST #2 PART 1 HAS BEEN PASSED!!!

[Part 2 - Stress test for a large number random ACEs for random files in export directory on the NFSv4 server] :

File: /nfs/nfs\_file1 / has been created

File: /nfs/nfs\_file1 / with errors

touch: cannot touch '/nfs/nfs\_file1': No such file or directory

File: /nfs/nfs\_file2 / has been created

File: /nfs/nfs\_file2 / with errors

touch: cannot touch '/nfs/nfs\_file2': No such file or directory

...

File: /nfs/nfs\_file20 / has been created

File: /nfs/nfs\_file20 / with errors

touch: cannot touch '/nfs/nfs\_file20': No such file or directory

ACE # 1 : setfacl -x u:nfs\_user14 /nfs/nfs\_file12

setfacl: /nfs/nfs\_file12: No such file or directory

...

ACE # 498 : setfacl -x u:nfs\_user16 /nfs/nfs\_file3

setfacl: /nfs/nfs\_file3: No such file or directory

ACE # 499 : setfacl -m g:nfs\_group51:xw /nfs/nfs\_file4

setfacl: /nfs/nfs\_file4: No such file or directory

ACE # 500 : setfacl -m g:nfs\_group14:xw /nfs/nfs\_file2

setfacl: /nfs/nfs\_file2: No such file or directory

[Check the status of a NFS service on the NFSv4 server after stress test] :

● nfs-server.service - NFS server and services

Loaded: loaded (/usr/lib/systemd/system/nfs-server.service; enabled; vendor preset: disabled)

Active: active (exited) since Thu 2016-06-02 18:46:57 MSK; 9min ago

Process: 994 ExecStart=/usr/sbin/rpc.nfsd \$RPCNFSDARGS (code=exited, status=0/SUCCESS)

Process: 993 ExecStartPre=/usr/sbin/exportfs -r (code=exited, status=0/SUCCESS)

Main PID: 994 (code=exited, status=0/SUCCESS)

CGroup: /system.slice/nfs-server.service

Jun 02 18:46:55 rhel systemd[1]: Starting NFS server and services...

Jun 02 18:46:57 rhel systemd[1]: Started NFS server and services.

THE TEST #2 PART 2 HAS BEEN PASSED!!!

[Clean created test data from the NFSv4 server] :

Created ACEs have been deleted from export dir and files in it on the NFS server.

Created files have been deleted from export dir on the NFS server.

Created users and groups have been deleted from the NFS server.

#####

Test #2 NFSv4 - Stress test for a large number of random operations

setting ACEs [Extended ACLs for UNIX] [PASSED]

[PASSED] [PASSED] [PASSED] [PASSED] [PASSED] [PASSED] [PASSED]

In order to get more information:

./logs/log\_run.log" - execution log (detailed information)

./logs/log\_result.log - log with the results

#####

## Test #3: NFSv4 - Stress test for a large number operations:

**[NFSv4 server side] : export / unexport**

**[NFSv4 client side] : mount / unmount**

=====

Test #3: NFSv4 - Stress test for a large number operations:

[NFSv4 server side] : export / unexport

[NFSv4 client side] : mount / unmount

Part 1 - Emulation of export a large number of directories on the  
NFSv4 server and their subsequent mounting on the NFSv4  
client, health check NFSv4 service

Part 2 - Emulation of a large number of random operations with  
exports on NFSv4 server and mounts on NFSv4 client,  
health check NFSv4 service

\*\*\* The test data will be created in the directory /mnt/\* from  
both sides NFSv4 client and server

=====

In order to run the test enter the required data:

1) Hostname of the NFSv4 server:

[input] : rhel

hostname have resolved

NFSv4 server: rhel

2) The number of exported (mounted) directories:

The number of directories [500..10000] [input] : 500

3) The number of random operations: export/unexport (server), mount/unmount (client):

The number of random operations [1000..20000] [input] : 1000

Great! The input data have been received!

After 5 seconds the test will be started...

=====

Client : ('Linux 4.4.9-300.fc23.x86\_64 x86\_64\n', None)

Hostname : fedora

IP : 192.168.100.182

Server : ('Linux 3.10.0-327.18.2.el7.x86\_64 x86\_64\n', None)

Hostname : rhel

IP : 192.168.100.176

Ping : fedora --->>> rhel

PING rhel (192.168.100.176) 56(84) bytes of data.

64 bytes from rhel (192.168.100.176): icmp\_seq=1 ttl=64 time=0.374 ms

64 bytes from rhel (192.168.100.176): icmp\_seq=2 ttl=64 time=0.395 ms

64 bytes from rhel (192.168.100.176): icmp\_seq=3 ttl=64 time=0.313 ms

--- rhel ping statistics ---

3 packets transmitted, 3 received, 0% packet loss, time 1999ms

rtt min/avg/max/mdev = 0.313/0.360/0.395/0.041 ms

[Part 1 - Stress test for export a large number of directories on the NFSv4 server  
and their subsequent mounting on the NFSv4 client] :

[Create 500 directories (both), export them (server) and mount them (client)] :

#1 STEP

[Server] : directory /mnt/nfs\_exp1 has been created

[Server] : directory /mnt/nfs\_exp1 has been exported

[Client] : directory /mnt/nfs\_mnt1 has been created

[Client] : directory /mnt/nfs\_exp1 has been mounted to local /mnt/nfs\_mnt1

...

#### #499 STEP

[Server] : directory /mnt/nfs\_exp499 has been created

[Server] : directory /mnt/nfs\_exp499 created errors

poll: protocol failure in circuit setup

[Server] : directory /mnt/nfs\_exp499 has been exported

[Server] : directory /mnt/nfs\_exp499 exported errors

poll: protocol failure in circuit setup

[Client] : directory /mnt/nfs\_mnt499 has been created

[Client] : directory /mnt/nfs\_exp499 has been mounted to local /mnt/nfs\_mnt499

[Client] : directory /mnt/nfs\_exp499 mounted errors

mount.nfs4: access denied by server while mounting rhel:/mnt/nfs\_exp499

#### #500 STEP

[Server] : directory /mnt/nfs\_exp500 has been created

[Server] : directory /mnt/nfs\_exp500 created errors

poll: protocol failure in circuit setup

[Server] : directory /mnt/nfs\_exp500 has been exported

[Server] : directory /mnt/nfs\_exp500 exported errors

poll: protocol failure in circuit setup

[Client] : directory /mnt/nfs\_mnt500 has been created

[Client] : directory /mnt/nfs\_exp500 has been mounted to local /mnt/nfs\_mnt500

[Client] : directory /mnt/nfs\_exp500 mounted errors

mount.nfs4: access denied by server while mounting rhel:/mnt/nfs\_exp500

[Check the final export directory on the server] :

[Check the final mount directory on the client] :

**THE TEST #3 PART 1 HAS BEEN FAILED!!!**

[Part 2 - Stress test for large number random operations: exports:

export/unexport (NFSv4 server), mount/unmount (NFSv4 client)] :

[Execute random operations with 1000 exports: export/unexport (server), mount/unmount (client)] :

#1 STEP [Server] : directory /mnt/nfs\_exp1 has been exported

#2 STEP [Server] : directory /mnt/nfs\_exp2 has been unexported

...

#997 STEP [Server] : directory /mnt/nfs\_exp997 has been exported  
 #997 STEP [Server] : directory /mnt/nfs\_exp997 exported errors  
 exportfs: Failed to stat /mnt/nfs\_exp997: No such file or directory  
 #998 STEP [Client] : directory /mnt/nfs\_exp998 has been mounted to local /mnt/nfs\_mnt998  
 #998 STEP [Client] : directory /mnt/nfs\_exp998 mounted errors  
 mount.nfs4: mount point /mnt/nfs\_mnt998 does not exist  
 #999 STEP [Client] : directory /mnt/nfs\_mnt999 has been unmounted  
 #999 STEP [Client] : directory /mnt/nfs\_mnt999 unmounted errors  
 umount: /mnt/nfs\_mnt999: mountpoint not found  
 #1000 STEP [Server] : directory /mnt/nfs\_exp1000 has been exported  
 #1000 STEP [Server] : directory /mnt/nfs\_exp1000 exported errors  
 exportfs: Failed to stat /mnt/nfs\_exp1000: No such file or directory

[Check the status of a NFS service on the NFSv4 server after stress test] :

- nfs-server.service - NFS server and services

Loaded: loaded (/usr/lib/systemd/system/nfs-server.service; enabled; vendor preset: disabled)

Active: active (exited) since Thu 2016-06-02 18:46:57 MSK; 23min ago

Process: 994 ExecStart=/usr/sbin/rpc.nfsd \$RPCNFSDARGS (code=exited, status=0/SUCCESS)

Process: 993 ExecStartPre=/usr/sbin/exportfs -r (code=exited, status=0/SUCCESS)

Main PID: 994 (code=exited, status=0/SUCCESS)

CGroup: /system.slice/nfs-server.service

Jun 02 18:46:55 rhel systemd[1]: Starting NFS server and services...

Jun 02 18:46:57 rhel systemd[1]: Started NFS server and services.

THE TEST #3 PART 2 HAS BEEN PASSED!!!

[Clean created test data from the NFSv4 client and server] :

[Client] : Mounted directories have been unmounted on NFSv4 client.

umount: /mnt/nfs\_mnt100: not mounted

...

umount: /mnt/nfs\_mnt9: not mounted

umount: /mnt/nfs\_mnt97: not mounted

[Client] : Old mounted directories have been deleted from NFSv4 client.

[Server] : Exported directories have been unexported on the NFSv4 server.

exportfs: Failed to stat /nfs: No such file or directory

[Client] : Old exported directories have been deleted from NFSv4 server.

#####

Test #3 NFSv4 - Stress test for a large number operations:

[NFSv4 server side] : export / unexport

[NFSv4 client side] : mount / unmount

[PASSED] [PASSED] [PASSED] [PASSED] [PASSED] [PASSED] [PASSED]

In order to get more information:

./logs/log\_run.log" - execution log (detailed information)

./logs/log\_result.log - log with the results

#####



# Test #4: The complex test for owner/permission/content modification of NFSv4 file system [POSIX]

Test #4: The complex test for owner/permission/content modification of  
NFSv4 file system [POSIX]

Part 1 - Check the ability of check the permissions of the files

- > Description : Check the ability of changes the permissions of the files
- > Steps : Run used functions of full\_generator
- > Expected result : Display "TEST #4 PART 2 HAS BEEN PASSED"

Part 2 - Check the ability of check the permissions of the directories

- > Description : Check the ability of changes the permissions of the directories
- > Steps : Run used functions of full\_generator
- > Expected result : Display "TEST #4 PART 2 HAS BEEN PASSED"

\*\*\* All parts of the complex test will be performed sequentially  
and will be running in the mount directory on NFSv4 client side

=====

In order to run the test enter the required data:

1) Hostname of the NFSv4 server:

[input] : rhel  
hostname have resolved  
NFSv4 server: rhel

2) Path to the exported (test) directory on NFSv4 server (/nfsdir, /mnt/nfs, ...):

[input] : /nfs  
NFS server's exported directory is existing  
NFSv4 exported dir: /nfs

3) The number files, directories to be created on the NFSv4 client in mount dir:

The number of objects [5..50] [input] : 10

Great! The input data have been received!

After 5 seconds the test will be started...

=====

Client : ('Linux 4.4.9-300.fc23.x86\_64 x86\_64\n', None)

Hostname : fedora

IP : 192.168.100.182

Server : ('Linux 3.10.0-327.18.2.el7.x86\_64 x86\_64\n', None)

Hostname : rhel

IP : 192.168.100.176

Ping : fedora --->>> rhel

PING rhel (192.168.100.176) 56(84) bytes of data.

64 bytes from rhel (192.168.100.176): icmp\_seq=1 ttl=64 time=0.271 ms

64 bytes from rhel (192.168.100.176): icmp\_seq=2 ttl=64 time=0.300 ms

64 bytes from rhel (192.168.100.176): icmp\_seq=3 ttl=64 time=0.310 ms

--- rhel ping statistics ---

3 packets transmitted, 3 received, 0% packet loss, time 1999ms

rtt min/avg/max/mdev = 0.271/0.293/0.310/0.025 ms

[Create 10 files, directories (random chmod) on the NFSv4 client in mount dir] :

File: /nfs/nfs\_file1 / has been created with permissions: u=w,g=wr,o=xr

File: /nfs/nfs\_file2 / has been created with permissions: u=xr,g=xwr,o=x

File: /nfs/nfs\_file3 / has been created with permissions: u=xw,g=w,o=xw

File: /nfs/nfs\_file4 / has been created with permissions: u=xw,g=x,o=xw

File: /nfs/nfs\_file5 / has been created with permissions: u=xr,g=r,o=x

File: /nfs/nfs\_file6 / has been created with permissions: u=w,g=xr,o=xwr

File: /nfs/nfs\_file7 / has been created with permissions: u=w,g=r,o=r

File: /nfs/nfs\_file8 / has been created with permissions: u=xwr,g=xw,o=xwr

File: /nfs/nfs\_file9 / has been created with permissions: u=xw,g=r,o=xwr

File: /nfs/nfs\_file10 / has been created with permissions: u=r,g=x,o=wr

Directory: /nfs/nfs\_dir1 / has been created with permissions: u=w,g=xwr,o=r

Directory: /nfs/nfs\_dir2 / has been created with permissions: u=r,g=wr,o=xwr

Directory: /nfs/nfs\_dir3 / has been created with permissions: u=xr,g=xw,o=x

Directory: /nfs/nfs\_dir4 / has been created with permissions: u=xr,g=x,o=x

Directory: /nfs/nfs\_dir5 / has been created with permissions: u=xwr,g=xwr,o=r

Directory: /nfs/nfs\_dir6 / has been created with permissions: u=xwr,g=xw,o=r

Directory: /nfs/nfs\_dir7 / has been created with permissions: u=xw,g=xwr,o=xwr

Directory: /nfs/nfs\_dir8 / has been created with permissions: u=xwr,g=x,o=xwr

Directory: /nfs/nfs\_dir9 / has been created with permissions: u=w,g=x,o=xwr

Directory: /nfs/nfs\_dir10 / has been created with permissions: u=xwr,g=xr,o=r

[Part 1 - Check the ability of check the permissions of the files] :

[Get directories and files in test directory /nfs] :

Get dir: /nfs/nfs\_dir\_444 permissions: 444

Get dir: /nfs/nfs\_dir\_666 permissions: 666

Get dir: /nfs/nfs\_dir1 permissions: 274

Get dir: /nfs/nfs\_dir2 permissions: 467

Get dir: /nfs/nfs\_dir3 permissions: 531

Get dir: /nfs/nfs\_dir4 permissions: 511

Get dir: /nfs/nfs\_dir5 permissions: 774

Get dir: /nfs/nfs\_dir6 permissions: 734

Get dir: /nfs/nfs\_dir7 permissions: 377

Get dir: /nfs/nfs\_dir8 permissions: 717

Get dir: /nfs/nfs\_dir9 permissions: 217

Get dir: /nfs/nfs\_dir10 permissions: 754

Get file: /nfs/nfs\_file\_444 permissions: 444

Get file: /nfs/nfs\_file\_666 permissions: 666

Get file: /nfs/nfs\_file1 permissions: 265

Get file: /nfs/nfs\_file2 permissions: 571

Get file: /nfs/nfs\_file3 permissions: 323

Get file: /nfs/nfs\_file4 permissions: 313  
Get file: /nfs/nfs\_file5 permissions: 541  
Get file: /nfs/nfs\_file6 permissions: 257  
Get file: /nfs/nfs\_file7 permissions: 244  
Get file: /nfs/nfs\_file8 permissions: 737  
Get file: /nfs/nfs\_file9 permissions: 347  
Get file: /nfs/nfs\_file10 permissions: 416

[Check the permissions of the files in test directory /nfs ] :

[Test to read for files with permissions: r] / Expected: PASSED

Try reading the file /nfs/nfs\_file\_444 / Expected: True  
Try reading the file /nfs/nfs\_file\_444 / Result: True  
Try writing <NFSv4 test> to the file /nfs/nfs\_file\_444 / Expected: False  
Try writing <NFSv4 test> to the file /nfs/nfs\_file\_444 / Result: True

[Test to read for files with permissions: r] / Result: PASSED

[Test to read and write for files with permissions: rw] / Expected: PASSED

Try reading the file /nfs/nfs\_file\_666 / Expected: True  
Try reading the file /nfs/nfs\_file\_666 / Result: True  
Try writing <NFSv4 test> to the file /nfs/nfs\_file\_666 / Expected: True  
Try writing <NFSv4 test> to the file /nfs/nfs\_file\_666 / Result: True  
Try reading the file /nfs/nfs\_file8 / Expected: True  
Try reading the file /nfs/nfs\_file8 / Result: True  
Try writing <NFSv4 test> to the file /nfs/nfs\_file8 / Expected: True  
Try writing <NFSv4 test> to the file /nfs/nfs\_file8 / Result: True

[Test to read and write for files with permissions: rw] / Result: PASSED

THE TEST #4 PART 1 HAS BEEN PASSED!!!

[Part 2 - Check the ability of check the permissions of the directories] :

[Get directories and files in test directory /nfs] :

Get dir: /nfs/nfs\_dir\_444 permissions: 444

Get dir: /nfs/nfs\_dir\_666 permissions: 666

Get dir: /nfs/nfs\_dir1 permissions: 274

Get dir: /nfs/nfs\_dir2 permissions: 467

Get dir: /nfs/nfs\_dir3 permissions: 531

Get dir: /nfs/nfs\_dir4 permissions: 511

Get dir: /nfs/nfs\_dir5 permissions: 774

Get dir: /nfs/nfs\_dir6 permissions: 734

Get dir: /nfs/nfs\_dir7 permissions: 377

Get dir: /nfs/nfs\_dir8 permissions: 717

Get dir: /nfs/nfs\_dir9 permissions: 217

Get dir: /nfs/nfs\_dir10 permissions: 754

Get file: /nfs/nfs\_file\_444 permissions: 444

Get file: /nfs/nfs\_file\_666 permissions: 666

Get file: /nfs/nfs\_file1 permissions: 265

Get file: /nfs/nfs\_file2 permissions: 571

Get file: /nfs/nfs\_file3 permissions: 323

Get file: /nfs/nfs\_file4 permissions: 313

Get file: /nfs/nfs\_file5 permissions: 541

Get file: /nfs/nfs\_file6 permissions: 257

Get file: /nfs/nfs\_file7 permissions: 244

Get file: /nfs/nfs\_file8 permissions: 737

Get file: /nfs/nfs\_file9 permissions: 347

Get file: /nfs/nfs\_file10 permissions: 416

[Check the permissions of the dirs in test directory /nfs ] :

[Test to read for dirs with permissions: r] / Expected: PASSED

Try reading the dir /nfs/nfs\_dir\_444 / Expected: True

Try reading the dir /nfs/nfs\_dir\_444 / Result: True

Try reading the dir /nfs/nfs\_dir\_444 / Expected: True

Try reading the dir /nfs/nfs\_dir\_444 / Result: True

[Test to read for dirs with permissions: r] / Result: PASSED

THE TEST #4 PART 2 HAS BEEN PASSED!!!

[Clean created test data from the NFSv4 client]

[Created test data have been cleaned from the NFSv4 client]

[Get directories and files in test directory /nfs] :

[Check the permissions of the dirs in test directory /nfs ] :

[Test to read for dirs with permissions: r] / Expected: PASSED

Try reading the dir /nfs/nfs\_dir\_444 / Expected: True

Try reading the dir /nfs/nfs\_dir\_444 / Result: False

Try reading the dir /nfs/nfs\_dir\_444 / Expected: True

Try reading the dir /nfs/nfs\_dir\_444 / Result: False

[Test to read for dirs with permissions: r] / Result: FAILED

THE TEST #4 PART 2 HAS BEEN FAILED!!!

#####

Test #4 The complex test for owner/permission/content modification of  
NFSv4 file system

[FAILED] [FAILED] [FAILED] [FAILED] [FAILED] [FAILED] [FAILED]

In order to get more information:

./logs/log\_run.log" - execution log (detailed information)

./logs/log\_result.log - log with the results

#####

## Test #5: Stress test for a large number of random operations setting

ACEs (Access Control Entries) [NFSv4 ACLs]

**Coming soon!!!**



